

VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning

SUBMITTED BY

1. **Rajdeep Das**
2. **Arjun Ghoshal**
3. **Arunava Kundu**

*Thesis submitted for the partial fulfillment of
the requirements for the degree
of*

BACHELOR OF TECHNOLOGY



(College Logo)



(University Logo)

**DEPARTMENT OF COMPUTER SCIENCE & BUSINESS SYSTEMS
INSTITUTE OF ENGINEERING & MANAGEMENT
MAULANA ABUL KALAM AZAD UNIVERSITY OF TECHNOLOGY,
WEST BENGAL
(Year) 2025**

VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning



A thesis submitted by

Rajdeep Das **22022002018001**

Arjun Ghoshal **12021002011004**

Arunava Kundu **22022002018005**

Supervisor 1

Prof.(Dr.) Indrajit De
*Dept. of Computer Science
& Business Systems
Institute of Engineering and
Management, Kolkata India*

Supervisor 2

Dr. Amitava Nag
*Dean of Academic Affairs
Central Institute of Technology,
Kokrajhar, India*

**DEPARTMENT OF COMPUTER SCIENCE & BUSINESS
SYSTEMS
INSTITUTE OF ENGINEERING & MANAGEMENT,
KOLKATA**

January 2025

DEPARTMENT OF COMPUTER SCIENCE & BUSINESS SYSTEMS



CERTIFICATE

This is to certify that the "**Thesis Report** on VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning " is submitted in partial fulfillment of the requirements for the degree of Bachelor of Technology by the following students:

Rajdeep Das	22022002018001
Arjun Ghoshal	12021002011004
Arnava Kundu	22022002018005

Prof. (Dr.) Indrajit De
*Dept. of Computer Science
& Business Systems
Institute of Engineering and
Management, Kolkata India*

Prof. (Dr.) Amitava Nag
*Dean of Academic Affairs
Central Institute of Technology,
Kokrajhar, India*

Head of the Department
*Department of Computer Science
& Business Systems*

Principal
IEM, Kolkata

Thesis Approval for B.Tech.

This thesis entitled "VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning" by Rajdeep Das is approved for the degree of Bachelor of Technology.

Examiner(s)

1.....

2.....

Date:

Place:

ACKNOWLEDGEMENT

I would like to express my deepest gratitude to all the individuals and institutions whose support and guidance were instrumental in the successful completion of this thesis.

First and foremost, I extend my sincere thanks to the Institute of Engineering and Management (IEM), Kolkata, for providing the academic environment, technical facilities, and mentorship that laid the foundation for this research. I am especially grateful to Prof.(Dr.) Indrajit De of the Department of Computer Science and Business Systems for his constant encouragement and invaluable feedback.

I am truly thankful to Dr. Amitava Nag and the Central Institute of Technology (CIT), Kokrajhar, for their technical support, expert insights, and research collaboration, which significantly contributed to the development and experimentation phases of this project.

My heartfelt appreciation also goes to the India Internet Foundation (IIFON) for providing essential infrastructure support that enabled smooth deployment, data communication, and testing within a simulated industrial network environment.

.....

(Signature)

.....

(Rajdeep Das, 22022002018001(54))

Date:

DECLARATION OF ORIGINALITY AND COMPLIANCE OF ACADEMIC ETHICS

I hereby declare that this thesis, VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning contains a literature survey and original research work carried out by me, the undersigned candidate, as part of my studies in the Department of Computer Science and Business System.

All information in this document has been obtained and presented in accordance with academic rules and ethical conduct. I also declare that, as required by these rules and regulations, I have fully cited and referenced all material and results that are not original to this work.

Details:

- Name: Rajdeep Das
- Examination Roll No: 22022002018001
- Registration No: 221040120854

I affirm that the work presented is original and all sources have been duly acknowledged.

Signature:

TABLE OF CONTENTS

Article No.	Content	Page No
1.	INTRODUCTION.....	3
1.1	WHAT IS IDS?	3
1.2	BACKGROUND.....	3
1.3	PROBLEM STATEMENT	6
1.4	AIMS AND OBJECTIVES.....	7
1.5	SIGNIFICANCE OF THE STUDY	8
1.6	SCOPE AND LIMITATIONS.....	9
2.	LITERATURE REVIEW	11
2.1	Federated & Transfer Learning-Based IDS Approaches	11
2.2	Privacy preserving and Quantum-Resilient Enhancements	12
3.	METHODOLOGY	17
3.1	Global Model Initialization	17
3.2	Local Fine-Tuning at Edge.....	17
3.3	Privacy Budget Allocation	17
3.4	Local Noise Injection	22
3.5	Robust Aggregation.....	23
3.6	Iterative Model Updates	27
3.7	Quantum Resilient Aggregation.....	27
4.	DEPLOYMENT.....	31
4.1	Front-End Deployment.....	31
4.2	Back-End Deployment	34
4.3	Network Deployment	51
5.	RESULTS AND ANALYSIS	59
5.1	Byzantine-Resilient Aggregation	59
5.2	Quantum-Resilient Communication.....	63

5.3	FRL+Nash Bargaining	65
5.4	Quantitative Improvements in System Analytics & Performance	67
6.	DISCUSSION AND CONCLUSION.....	69
6.1	Discussion	69
6.2	Conclusion.....	70
7.	SUMMARY, PUBLICATIONS AND FUTURE WORK.....	72
7.1	FUTUTE WORK	73
7.2	PUBLICATIONS	74
	APPENDIX-A: PROTOTYPE SNAPS.....	75
	REFERENCES.....	79

ABSTRACT

Traditional Intrusion Detection Systems (IDS) frequently fail to scale, adapt, and protect data privacy in the ever-changing world of cyber threats, particularly in Enterprise-level networks environments. In order to provide safe, decentralised, and privacy-preserving anomaly detection, this thesis presents VANGUARD, a next-generation IDS framework that makes use of Federated Transfer Learning (FTL), Differential Privacy (DP), and Quantum-Resilient Communication. A Random Forest model that has been globally trained on the CIC-IDS-Collection dataset serves as the foundation for the suggested architecture. Under stringent DP constraints enforced by Gaussian noise injection, this model is distributed to local edge nodes, where local data is used for fine-tuning. By combining Nash Bargaining with a Reinforcement Learning (RL)-driven mechanism, privacy budgets are dynamically distributed to participating clients, guaranteeing the best possible trade-offs between privacy and utility. The return of locally updated models in encrypted batches is facilitated by a secure Kafka-based communication pipeline, where robust model integration is ensured by Multi-Krum Byzantine Fault-Tolerant aggregation, supplemented by reputation-based weighting. Kyber-based key encapsulation and AES-256-CBC encryption are used to strengthen communication for post-quantum resilience. VANGUARD protects sensitive data and upholds operational integrity while achieving real-time, adaptive intrusion detection. By using this framework, the study not only overcomes the shortcomings of the current IDS but also establishes a standard for safe and scalable threat detection systems in vital enterprise infrastructures.

Chapter-1

INTRODUCTION

1.1 WHAT IS IDS?

The intrusion detection system (IDS) operates as a security tool that examines network and system operations to track suspicious activity and unauthorized access and policy breaches. A cybersecurity framework heavily relies on this protective watchful system which observes user behavior and system logs or data packets during periodic or real-time scans to detect possible threats. The main distinction between IDS and firewalls shows that IDS identifies suspicious activities which bypass primary security systems but firewalls prevent unwanted access. The IDS alerts administrators through alerts so they can receive threat detection details including attack type alongside source IP addresses and destination IPs as well as compromised system information and threat evaluation level.

The two core intrusion detection systems include Network-based IDS (NIDS) and Host-based IDS (HIDS). HIDS tracks individual devices or hosts for their system anomalies plus unauthorized file modifications and incorrect system settings yet NIDS surveys network traffic to spot known attacks or abnormal network activities. IDS systems incorporate two detection methods including anomaly-based detection and signature-based detection. The intrusion detection system (IDS) functions as the essential component for proactive security measures because it provides better situational awareness as well as analytical capabilities for forensics while enabling swift response to security threats independently of its threat blocking capabilities.

1.2 BACKGROUND

Modern network environments require intelligent IDS which offer flexibility alongside privacy protection because of expanding cyber threats. IDS stands for Intrusion Detection System which functions as a cybersecurity instrument that helps track and analyze abnormal system activity to make detection possible. IDS functions to discover potential hazards through its examination of intrusions together with its search for data breaches and DDoS attacks as well as its surveillance for internal misuse. The system enables monitoring either during offline periods or real-time operations. Two principal detection methods exist within IDS solutions which combine anomaly detection with signature detection. Signature-based systems analyze incoming traffic by comparing it to known attack patterns for effectiveness yet remain useless against zero-day threats.

Anomaly-based intrusion detection systems, on the other hand, use statistical models or machine learning to identify departures from known normal behaviour. This enables them to detect threats that have not yet been identified, but at the risk of more false positives. The current industry demands IDS solutions as a fundamental protection mechanism for critical infrastructure alongside enterprise systems and both industrial IoT (IIoT) and cloud platforms. These systems create the first barrier of security protection which enhances the capabilities of firewalls and endpoint protection tools. Security information and event management solutions have spread across finance institutions as well as healthcare facilities and manufacturing organizations and governmental bodies as they protect their networks from unauthorized activities while maintaining data privacy and meeting regulatory criteria. When threats change and data spreads across numerous, private systems, traditional IDS techniques become ineffective. In complex and uncertain system environments, static rules, centralised data storage techniques, and unresponsive systems perform poorly. The current VANGUARD project is the result of a multi-year study conducted by our research team to examine the need for a shift to intelligent distributed privacy-aware IDS solutions due to the evolution of security threats.

We started with Athena, a forecasting technique that uses conventional Machine Learning (ML) algorithms to identify and lessen Distributed Denial of Service (DDoS) attacks. Traditional machine learning techniques with Athena demonstrated satisfactory performance in enhancing initial threat identification capabilities. The study presented at ICIS2024 in Hanoi Vietnam showcased weight reduction models to recognize threats in real-time while achieving recognition at this conference. Athena simultaneously exposed limitations related to growing operational needs and protection against concealed security risks as well as its dependence on centralized information storage systems.

We developed the reliable security testbed IoTForge Pro for simulating Industrial IoT (IIoT) environments to solve the issues with obtaining realistic intrusion datasets and IDS evaluation experiments[13]. The system established a method to develop multiple datasets that replicated real-world IIoT security attacks. This project produced both ForgeIIoT dataset in addition to modular test framework which was published in the IEEE Internet of Things Journal. This initiative established an extensive empirical basis used for IDS model training and validation inside authentic industrial settings [Figure 1.1].



Figure 1.1 IoTForge Pro Testbed Setup

Using previous foundations, we initiated the Aegis project which developed Federated Transfer Learning (FTL) with built-in Differential Privacy (DP). The first-time implementation of Gaussian noise within local model updates prevents sensitive client data from model inversion attacks. Secure decentralised model training became achievable through Aegis by implementing the Federated Transfer Learning (FTL) system with FedAvg. The innovative approach received approval at WIN6.0 during the gathering held in Kolkata India.

We envisioned FORTRESS, a next-generation IDS designed for Industrial IoT that integrates FTL, Blockchain, and Post-Quantum Cryptography (PQC), as an extension of Aegis. In order to guarantee operational security even in hostile and resource-constrained environments, the project sought to completely decentralise the IDS pipeline. FORTRESS enhanced our federated model-sharing mechanism with tamper-proof recordkeeping and blockchain-backed transparency. At ACM CCS 2025, the project is presently being reviewed.

As a result of these combined efforts, we were able to determine the need for a more flexible, secure, and reliable method that combines adversarial resilience, federated learning, differential privacy, and intelligent resource negotiation. Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning, or VANGUARD, is a comprehensive framework that incorporates all of our earlier innovations while addressing their shortcomings. VANGUARD, which embodies cutting-edge techniques for federated, privacy-preserving, and post-quantum secure intrusion detection in dynamic, distributed networks, is the pinnacle of our research development.

1.3 PROBLEM STATEMENT

IDS systems face major hurdles in detecting threats effectively because of advanced cyber threats that exist in diverse network environments such as IoT, IIoT, enterprise-level and critical service platforms. The common utilization of static detection techniques within centralized architectural designs of intrusion detection systems hinders effective security of sensitive data while preventing quick adaptation to evolving attack patterns and efficient operation across multiple locations. Sensitive user or organisational data is also exposed to potential privacy risks, legal infractions, and model inversion attacks because the majority of centralised IDS approaches require the collection of raw data at a central server. Federated learning (FL), which allows for the decentralised training of IDS models without the sharing of raw data, presents a promising substitute; however, it is still susceptible to dangers such as gradient leakage, poisoned updates, and adversarial interference during aggregation.

System weaknesses stem from the lack of secure communication protocols uniting separate IDS nodes. The consistent scalability of FL-based IDS frameworks struggles during real-world deployments since they lack adequate tools for privacy budget negotiation and trust-aware aggregation and adversarial manipulation resistance. Most existing solutions fail to establish proper system-wide update protocols which maintain fresh models and sustained protection abilities across client networks. A major threat to infrastructure security and data protection develops because IDS communication pipelines do not use post-quantum cryptography while current systems maintain dangerous weaknesses to future cryptanalysis attacks. The current situation demands an immediate development of a complete intrusion detection framework that features scalability and future-readiness:

- Enables **decentralized training with differential privacy guarantees**,
- Incorporates **reinforcement learning-based privacy budgeting**,
- Utilizes **reputation-aware, adversary-tolerant aggregation strategies**,
- Supports **continuous and iterative model refinement**, and
- Ensures **post-quantum secured communication** across nodes.

Our proposed system, **VANGUARD: Versatile Aggregation for Networked Governance using Unified AI with Reinforced Decisioning**, directly addresses these limitations, offering a resilient, adaptive, and privacy-preserving IDS architecture suitable for modern and future digital ecosystems.

1.4 AIMS AND OBJECTIVES

Designing and developing VANGUARD, a novel, distributed, and privacy-preserving intrusion detection system (IDS), is the main goal of this project. It uses post-quantum cryptography, differential privacy, and advanced machine learning techniques to provide secure, adaptive, and reliable threat detection in contemporary networked environments.

The following goals serve as a guide for the project in order to achieve this goal:

→ **Design a Federated Transfer Learning (FTL) Framework**

Develop a scalable FTL-based architecture for intrusion detection, where a globally pre-trained model is distributed to edge nodes and fine-tuned on locally collected data, thereby eliminating the need for raw data sharing.

→ **Integrate Differential Privacy using Reinforcement Learning and Nash Bargaining**

Introduce a Reinforcement Learning (RL) and Nash Bargaining-based privacy budget allocation mechanism that dynamically controls Gaussian noise addition to local datasets, preserving privacy and resisting model inversion attacks.

→ **Implement Trust-Aware Byzantine-Resilient Aggregation**

Employ a Byzantine Fault Tolerant (BFT) strategy — Multi-Krum — combined with a reputation-based weighting system to ensure secure, adversary-resilient aggregation of fine-tuned models at the server node.

→ **Enable Seamless and Iterative Model Updates**

Design a cyclic update mechanism where each newly aggregated global model is merged with the previous one, maintaining continuous improvement and freshness in deployed IDS agents across client systems.

→ **Ensure Secure, Future-Proof Communication**

Secure all inter-node and client-server communications using AES-256-CBC encryption, with session keys encapsulated via **Kyber** (a post-quantum secure Key Encapsulation Mechanism), protecting the system against classical and quantum cryptographic threats.

→ **Deploy High-Throughput, Real-Time Data Pipelines**

Utilizing Apache Kafka, Redis Timeseries and Webhooks to facilitate high-speed, fault-tolerant communication between distributed nodes, enabling real-time streaming of model updates and attack detection signals.

→ **Validate Performance using State-of-the-Art Datasets**

Train and evaluate the system using the **CIC-IDS-Collection**, a comprehensive suite of benchmark datasets covering a wide range of modern cyberattacks, to assess accuracy, scalability, privacy preservation, and resilience under attack.

Through the above objectives, **VANGUARD** aspires to contribute a future-ready IDS framework capable of defending next-generation digital infrastructures from ever-evolving cybersecurity threats.

1.5 SIGNIFICANCE OF THE STUDY

Commercial and ecosystem security depends heavily on next-generation intrusion detection systems because adversaries now employ sophisticated digitally evasive cyber threats against industries nationwide. **VANGUARD** establishes a significant breakthrough in intrusion detection by integrating Post-Quantum Cryptography with Federated Transfer Learning and Differential Privacy to develop a single integrated IDS framework.

Intrusion detection systems with conventional designs need processing to occur in one central location through continuous data transfer to their central server for detection or training purposes. The delivery of sensitive network data exposes it to problems with scalability and security threats as well as legal complications. The current methods fail to resist future quantum computing threats and they lack advanced protection against attacks through model inversion or poisoning.

VANGUARD directly addresses these limitations by offering:

- **Decentralized Learning:** Reduces reliance on central data repositories, minimizing privacy leakage and meeting data sovereignty requirements.
- **Differential Privacy Integration:** Protects the integrity of user data at the edge level before model training, using intelligent, game-theory-based budget allocation.
- **Byzantine Fault Tolerant Aggregation:** Ensures that adversarial nodes or poisoned model updates are identified and excluded, increasing system robustness.
- **Post-Quantum Security:** Implements encryption mechanisms like AES-CBC with Kyber encapsulation to future-proof communications against quantum threats.
- **Continuous Adaptation:** Through iterative update cycles, the IDS remains current with evolving threat landscapes, offering real-time, intelligent defense.

Real-time operations and industrial implementation in ELN(enterprise-level networks) become possible through the integration of Apache Kafka and fast data transfer pipelines.

Through this research the authors build fundamental elements for future IDS systems which combine scalability with intelligence and secure-by-design competence. The proposed framework serves academia and industry with a fresh approach while underlining the importance of privacy-protecting decentralized cyber protection solutions for current global security risks.

1.6 SCOPE AND LIMITATIONS

1.6.1 SCOPE

The development and assessment of VANGUARD, a post-quantum secure, scalable, and privacy-preserving intrusion detection system (IDS) built on Federated Transfer Learning (FTL), is the main focus of this project. With a focus on applications in ENets, critical infrastructure, multi-tenant cloud environments, and Industrial IoT (IIoT), the system is built to operate in a distributed network environment.

The key components within the scope include:

- **Federated Training:** A global Random Forest model is trained on CIC-IDS-Collection and distributed to local edge nodes for fine-tuning with locally relevant datasets.
- **Privacy Preservation:** Integration of Differential Privacy with Reinforcement Learning and Nash Bargaining Theory to allocate privacy budgets and protect against model inversion attacks.
- **Adversarial Robustness:** Use of Byzantine-resilient aggregation (Multi-Krum) and reputation-based weighing to mitigate poisoning attacks during model updates.
- **Real-Time Model Streaming:** Use of Apache Kafka for asynchronous batched model update communication between edge and server.
- **Post-Quantum Cryptographic Security:** All communication channels secured using AES-256 CBC encryption with session key encapsulated using Kyber, offering resistance to quantum-era threats.
- **Continuous Learning Framework:** Iterative model refinement ensures that client-deployed IDS models remain updated and effective against emerging threats.

1.6.2 LIMITATIONS

Despite its powerful design and architectural creativity VANGUARD suffers from certain implementation issues.

- **Simulated Environment:** We relied on CIC-IDS-Collection and artificial edge frameworks as benchmark datasets to evaluate the model within laboratory testing conditions. The actual network deployment environment lacks robust modeling of heterogeneous hardware and erratic network conditions and real-time packet loss behavior.
- **Limited Attack Vectors:** The broad attack vector coverage of the dataset still leaves space for current zero-day threats to avoid detection because the dataset does not encompass them.
- **Resource Overheads:** Byzantine-resilient aggregation and differential privacy result in computational overhead. The resources on edge devices with restricted capabilities might experience both delayed operations and diminished performance because of these factors.
- **Model Specificity:** The Random Forest core model used in the system lacks simplicity and interpretability compared to shallow learners or rule-based systems which may hinder efficient domain generalization.
- **Security Assumptions:** The security model of the system works with a semi-honest threat model which implies honest-but-curious nodes yet it fails to protect against all types of advanced persistent threats or physical security attacks.

Chapter-2

LITERATURE REVIEW

Recent years have witnessed substantial advancements in Intrusion Detection Systems (IDS) development that achieves both effectiveness and privacy protection because of augmented cybersecurity threats alongside decentralized data. Thanks to the continuously increasing complexity and volume of attacks on ENets, researchers have started to use Federated Learning (FL) and Transfer Learning (TL) and Differential Privacy (DP) alongside Quantum-Resilient Security Models. The review analyzes scholarly works in addition to practical developments in the field and their features against the functionality and constraints of commercial IDS systems.

2.1 Federated and Transfer Learning-Based IDS Approaches

Zhang et al. put forward a federated learning based anomaly detection method that can be implemented in IIoT systems[1]. In their work, they have used instance-based transfer learning to address one of the characteristics of IIoT network, non-independent and non-identically distributed (non-IID) data. The system used rank aggregation with a voting scheme and demonstrated better results in detection than conventional benchmarks. Importantly, based on the proposed framework, classification accuracies of 95.97 %, and 73.70 % for AdaBoost, and Random Forest, respectively have been obtained indicating much benefits of integrating federated learning with transfer learning to enhance privacy-preserving and efficient intrusion detection in IIoT scenario.

Building on that, Otoum et al highlighted on how federated and transfer learning could be beneficial in the classification of intrusion for IoT devices[2]. In that work they also focused on the importance and benefits of using such approaches in enhancing model learning rates, reducing training data needs, and increasing model performance, all the while maintaining users' privacy intact. Through adopting the IoMT as a use case, the study established tremendous enhancements in the intrusion detection confirming the adaptability of FL and transfer learning applied on IoT applications of broad areas. Subsequently, federated learning frameworks have been further enriched by self-supervised learning (SSL).

Meyer et al. presented the Federated Self-Supervised Learning (FSSL) framework which combines the SSL concept with federated learning due to the lack of labeled data[3]. FSSL successfully improved the IDSs' performance by training autoencoders at the client level to learn data representations. This framework appeared to possess higher F1-scores over various datasets such as NSL-KDD, TonIoT, and BotIoT, and especially effective in discovering new types of attack. This last approach was not only beneficial for increasing the accuracy of the created models, but stressed the significance of the use of the data containing no labels for effective intrusion detection.

In a situation where there is limited availability of resources, Cheng's et al have presented a federated transfer learning framework that involves RL based client selection[4]. This method enhanced the efficiency of the clients within the federated learning to contribute into the global model such that only quality updates were incorporated. When incorporating RL with federated transfer learning, the framework enhanced the accuracy of intrusion detection and the efficiency of communication in applications in MEC context. This served to underscore the need for 'smart' client selection in light of issues of scalability and resource availability.

2.2 Privacy-Preserving and Quantum-Resilient Enhancements

Apart from federated and transfer learning, using quantum-resilient features and blockchain to enhance security has been seen effective in critical Infrastructures. Xu et al. put forward cross-layer authentication schemes based on quantum encryption to protect the device's conversation in the context of 5G empowered IIoT[5]. For encoding a device identity, their system integrated emitting quantum walks, propagating high privacy and scalability. The proposed framework was also resilient to both quantum and classical threats and will be a futuristic approach to IIoT authentication.

In the same way, Ghaemi and Abbasinezhad-Mood designed a novel self-certified authentication protocol involving blockchain and quantum computing immunity to cross industrial relations[6]. As their protocol worked out the flaws of the existing security models based on blockchain such as the intermediation of the servers and vulnerability to the quantum attacks. Requiring only self-certified authentication and implementing a decentralized ledger, the framework brought down communication and computation overheads by 13\% and 91\% compared to the existing solutions. This idea highlighted the concept of the utilization of the integration of blockchain and quantum-protected cryptography to counter the novel risks in IIoT environments.

Almesleh et al. designed an FL-DP model specifically intended for IoT devices with limited resources which offers privacy protection without performance loss. The implementation of DP mechanisms succeeded but highlighted that processing capability gets reduced when aiming for privacy protection[34].

Using a hybrid DNN-CNN architecture Rajesh et al. implemented FTL to process high-dimensional IIoT data. The system delivered solid detection results but experienced running limitations and encountered problems with redundant features. The PrivateKT method developed by the authors successfully optimized FL knowledge transfer by minimizing the difference between FL and centralized models operating in full DP mode[29].

Internal research at Meta involving mobile applications using FL-DP demonstrated the significance of resolving label imbalance and handling heterogeneous data distribution which present major challenges in operational federated systems.

2.3 Industrial-Grade IDS Tools: A Practical Landscape

Academic research leads in developing adaptive IDS systems with privacy measures but the cybersecurity sector provides multiple dominant tools across the marketplace. Modern research aims to resolve the important shortcomings that exist in contemporary solutions. Our discussion requires first an examination of modern IDS deployments.

2.3.1 Snort

Snort maintains its position as one of the most popular open-source IDS because Cisco manages it. This IDS detects threats through signature matching while providing comprehensive rule controls and considerable community backing and plugin distribution. The known threat detection abilities of Snort make it effective yet it fails to handle zero-day attacks alongside encrypted traffic which also produces centralized operating challenges for heterogeneous networks and fails to protect privacy of sensitive data. Snort finds its best fit in organizations that can maintain signature databases because it does not effectively manage APTs alone but it serves small to medium businesses as well as educational institutions alongside organizations that employ dedicated security teams to maintain signature databases. Pattern matching in this solution proves ineffective when defending against polymorphic malware and complex attacks that aim to avoid signature detection so security personnel need multiple defense elements while also maintaining active signature maintenance.

2.3.2 Zeek (formerly Bro)

Zeek operates as an advanced network monitoring framework that transforms network traffic to security events via its programming code to help users build dedicated detection patterns and event correlation models while employing protocol inspection for application-level detection. Zeek offers exceptional advantages through its scripting language but its research extendibility results in performance decline when dealing with complicated syntax which hampers extensive network expansion and only seasoned experts can successfully use it. Associations unable to retain professional security personnel will steer clear of Zeek implementation because of its complex learning requirements even though they would need to integrate host-based protection to achieve total security protection. Research laboratories and large businesses employing security teams for network investigation can utilize the solution because of its complex customization requirements and maintenance complexity.

2.3.3 Cisco Secure IDS

Enterprise organizations receive benefit from Cisco Secure IDS through commercial security solutions operating in Cisco security infrastructure which provides system performance enhancements and automatic threat detection from Cisco Talos Intelligence Group as well as native compatibility with Cisco security platforms. Enterprise subscribers use this security system in conventional administration hubs at the cost of subscription fees which blocks domain-specific customization options. The solution presents standard behavioral analysis operations without edge computing capabilities and leads to difficulties when integrating with security solutions outside Cisco's ecosystem. This solution performs optimally because of proper deployment and ongoing upkeep yet it requires experts in Cisco systems who ensure its capability reaches maximum effectiveness for extensive Cisco networks that seek vendor collaboration and extensive documentation.

2.3.4 Darktrace

Darktrace enables advanced next-generation AI functionality by using unsupervised machine learning anomaly detection to automatically learn organizational systems in combination with Entity Behavior Analytics that detects possible compromise indicators older systems tend to overlook.

The proprietary system of Darktrace offers limited transparency about decision processes and maintains high acquisition expenses which restricts the system from smaller organizations' reach. Organizations face difficulties in privacy compliance across multiple jurisdictions because this system utilizes a centralized artificial intelligence framework which also results in exorbitant security alerts unless organizations develop flawless alert investigation methods. The solution's operational effectiveness depends on network complexity during baseline establishment yet delivers maximum success for institutions focused on financial data security and critical infrastructure together with security centers capable of handling AI-generated alerts.

Below table shows a detailed comparison between our [roposed suite with the existing state-of-the-art suites and tools [Table 1.1]:

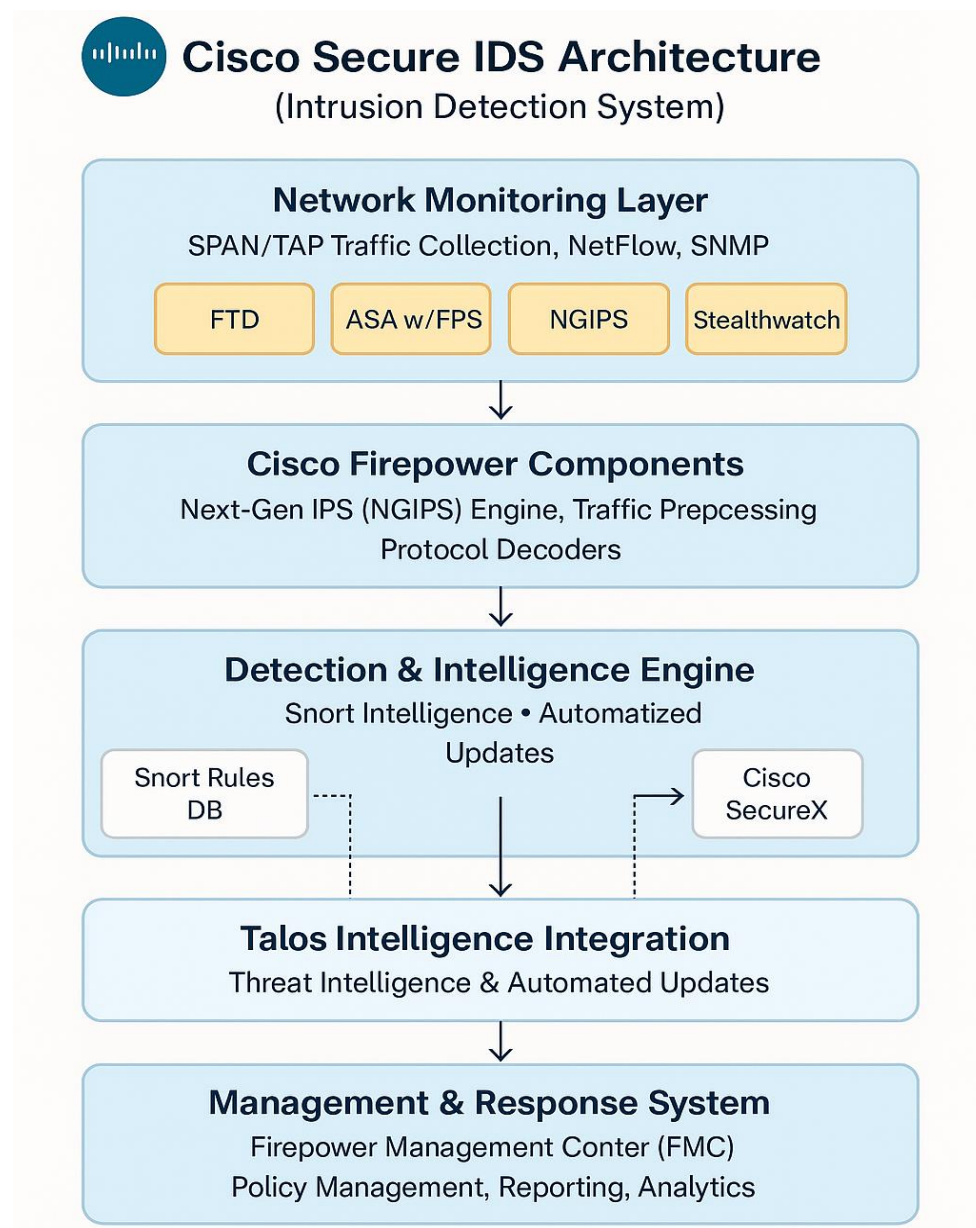


Figure 2.1 Cisco Secure IDS Workflow

Chapter-3

METHODOLOGY

Our proposed project, **VANGUARD** framework introduces a privacy-preserving, federated IDS architecture tailored for Enterprise Networks (Enets). It integrates byzantine resilient federated transfer learning, differential privacy with reinforced budget allocation, and quantum-resilient communication to ensure scalable, secure, and adaptive intrusion detection.

3.1 Global Model Initialization

The CIC-IDS-Collection, a comprehensive dataset made up of CICIDS2017, CIC-DoS2017, CSE-CIC-IDS2018, and CIC-DDoS2019, is used to train the first global model[7]. Numerous attack vector types and network traffic patterns are included in this varied dataset. For large-scale traffic classification, a Random Forest classifier is used because of its interpretability and resilience to overfitting. Baseline intrusion patterns from various attack classes, including DoS, brute force, and botnets, are captured by the global model. After this, the pre-trained Global model is sent to the main federated cloud server, that is, AWS EC2 server.

3.2 Local Fine-Tuning at Edge

From the server, the pretrained Global Model is being distributed among all the local networks in the edge layer for finetuning with the locally specific network data. This enables **domain-specific adaptation** through transfer learning while preserving decentralized data ownership.

3.3 Privacy Budget Allocation

3.3.1 Differential Privacy

Differential Privacy (DP) provides a formal guarantee about the maximum influence an individual's data can have on the output of an analysis. A randomized mechanism M satisfies ϵ -differential privacy if for any two adjacent datasets D and D' differing by at most one record, and for all possible outputs $S \subseteq \text{Range}(M)$:

$$\Pr[M(D) \in S] \leq e^\epsilon \cdot \Pr[M(D') \in S]$$

The parameter $\epsilon > 0$ is the privacy budget, which quantifies the privacy guarantee: smaller values of ϵ provide stronger privacy protection but typically at the cost of reduced utility.

In the context of federated learning, DP is commonly implemented by adding calibrated noise to client updates before they are shared with the server. The amount of noise is proportional to the sensitivity of the computation and inversely proportional to the privacy budget. For client k with privacy budget ϵ_k , the noisy update becomes:

$$\widetilde{\Delta w}_k^t = \Delta w_k^t + \mathcal{N}(0, \sigma_k^2 \cdot I)$$

where σ_k is the noise scale, determined by the sensitivity of the update operation and the desired privacy level ϵ_k .

3.3.2 Nash Bargaining

The Nash Bargaining Solution (NBS) addresses fair resource allocation in cooperative settings. For clients $k \in \{1, 2, \dots, N\}$ with utility functions U_k and disagreement points U_k^* (representing minimum acceptable utilities), the NBS maximizes the Nash product:

$$\max_{\{x_k\}} \prod_{k=1}^N (U_k(x_k) - U_k^*)$$

In our context, the resources being allocated are privacy budgets ϵ_k , with the constraint that:

$$\sum_{k=1}^N \epsilon_k \leq \epsilon_{\text{total}}$$

3.3.3 Reinforcement Learning

Reinforcement Learning (RL) involves an agent learning to make decisions by interacting with an environment. The process is modeled as a Markov Decision Process (MDP) defined by a tuple (S, A, P, R, γ) , where: - S is the state space - A is the action space - $P : S \times A \times S \rightarrow [0, 1]$ is the transition probability function - $R : S \times A \times S \rightarrow \mathbb{R}$ is the reward function - $\gamma \in [0, 1)$ is the discount factor. The agent's goal is to learn a policy $\pi : S \rightarrow A$ that maximizes the expected cumulative discounted reward:

$$\mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \right]$$

Policy optimization methods, such as Proximal Policy Optimization (PPO), iteratively improve the policy based on collected experiences. In our framework, each client implements an RL agent whose state incorporates local data characteristics and negotiation history, actions represent proposed privacy budget adjustments, and rewards reflect improvements in both individual utility and the Nash product.

3.3.4 PROPOSED FRL-BASED NASH BARGAINING ALGORITHM

The DP budget negotiation problem is formulated as a cooperative Nash Bargaining game. For each client $k \in \{1, \dots, N\}$, let $U_k(\epsilon_k)$ denote the utility derived from a privacy budget ϵ_k , and U_k^* denote the disagreement point (minimum acceptable utility). The objective is to maximize the Nash product:

$$\begin{aligned} & \max_{\{\epsilon_k\}} \prod_{k=1}^N (U_k(\epsilon_k) - U_k^*) \\ & \text{subject to} \\ & \sum_{k=1}^N \epsilon_k \leq \epsilon_{\text{total}}, \quad \epsilon_k > 0 \quad \forall k. \end{aligned}$$

Taking logarithms, we obtain a tractable formulation:

$$\max_{\{\epsilon_k\}} \sum_{k=1}^N \log (U_k(\epsilon_k) - U_k^*)$$

3.3.4.1 FRL Framework

Each client in our system maintains an individual reinforcement learning (RL) agent that continuously learns an optimal negotiation strategy. The RL agent interacts with its environment—defined by its local data, budget allocation, and negotiation history—to dynamically adjust its privacy budget allocation. The key components of this framework are:

- **State $\mathbf{s}_k^{(t)}$:** The state vector encodes all relevant information needed for decision-making at time step t . It typically includes:
 - ✓ The current privacy budget allocation $\epsilon^{(t)}_k$ for client k .
 - ✓ Historical negotiation outcomes, such as previous utility values or past adjustments.
 - ✓ Local data characteristics, for instance, data quality metrics, data size, and sensitivity measures computed from the client's local dataset.
 - ✓ Global negotiation parameters such as the total privacy budget ϵ_{total} or summary statistics aggregated from other clients.

This comprehensive state representation enables the RL agent to assess both its own performance and the overall system context, thus tailoring its negotiation actions accordingly.

- **Action $a^{(t)}_k$:** The action represents a proposed adjustment to the current privacy budget. Formally, it is sampled from the client's policy distribution:

$$a_k^{(t)} \sim \pi_{\theta_k} \left(s_k^{(t)} \right)$$

where π_{θ_k} is the policy parameterized by θ_k . The action is typically a continuous value that suggests an increment or decrement in the allocated budget. The design of the action space is critical: it must allow fine-grained adjustments so that the negotiation can converge smoothly while exploring different allocation strategies.

- **Reward $r^{(t)}_k$:** The reward quantifies the benefit of the agent's action by reflecting improvements in both the individual utility and the collective Nash bargaining objective. It is computed based on:
 - ✓ The change in the client's utility, typically measured in logarithmic terms to capture proportional gains.
 - ✓ The improvement in the overall Nash product, which ensures that the negotiation not only benefits individual clients but also enhances the joint outcome.

The reward signal drives the RL agent to adjust its strategy over time, promoting actions that yield higher utility and contribute to a more equitable distribution of the global privacy budget.

In summary, the Federated Reinforcement Learning framework allows each client to iteratively learn from its local environment and negotiation history, enabling adaptive and fair adjustments to privacy budget allocations across the entire federated system.

3.3.4.2 Negotiation Process

The FRL-based Nash bargaining negotiation proceeds iteratively as follows [Figure 3.1]:

- *Initialization:* The server first announces the global privacy budget ϵ_{total} . Each client then initializes its local budget as:

$$\epsilon_k^{(0)} = \frac{\epsilon_{\text{total}}}{N},$$

ensuring an equal starting point for all clients. This step sets the baseline from which subsequent adjustments will be made.

- *Local Computation:* Each client uses its local dataset to compute key metrics:

- The utility $U_k(\epsilon_k^{(t)})$, which quantifies the benefit of receiving a certain privacy budget allocation. This utility typically depends on factors like data quality and size.
- The sensitivity S_k , representing how responsive the client's output is to changes in the data.

Using these metrics, the client forms its state representation $s_k^{(t)}$, which encapsulates its current budget, utility, sensitivity, and potentially other features (e.g., historical performance).

→ *Action Selection*: Each client's RL agent observes its state $s_k^{(t)}$ and samples an action: $a_k^{(t)} \sim \pi_{\theta_k}(s_k^{(t)})$,

where π_{θ_k} is the client's policy parameterized by θ_k . The selected action corresponds to a proposed adjustment $\Delta\epsilon_k^{(t)}$, determined by a function $f(a_k^{(t)})$. This action reflects the client's strategy to either request a higher or lower share of the privacy budget based on its current state.

→ *Aggregation*: The server collects all proposed budget adjustments $\{\Delta\epsilon_k^{(t)}\}_{k=1}^N$ from the clients. It then updates each client's budget as:

$$\epsilon_k^{(t+1)} = \epsilon_k^{(t)} + \Delta\epsilon_k^{(t)},$$

and normalizes the updated budgets to ensure that the global constraint is met:

$$\sum_{k=1}^N \epsilon_k^{(t+1)} = \epsilon_{\text{total}}$$

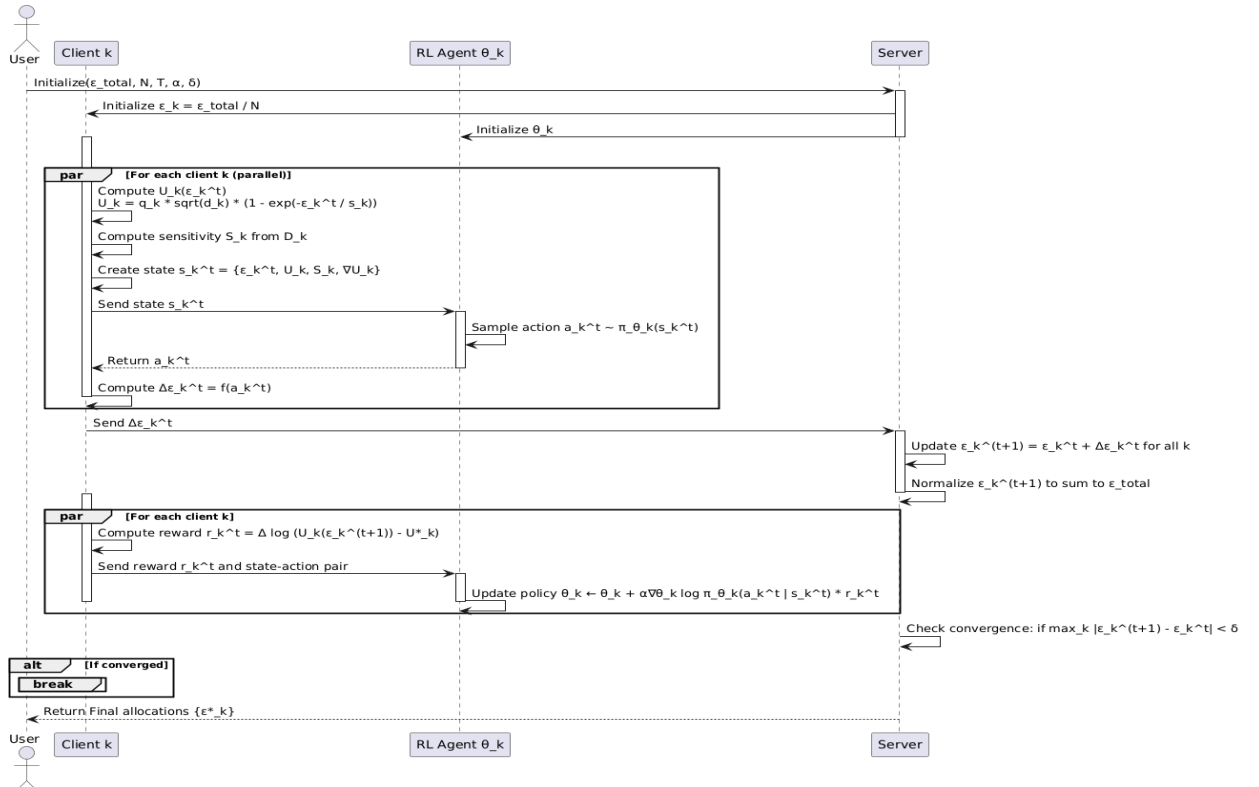


Figure 3.1 FRL-based Nash Bargaining for DP Budget Negotiation Workflow

3.4 Local Noise Injection

In order to preserve privacy-providing properties within the federated learning pipeline, Local Noise Injection is proposed after the FRL-based Nash bargaining and prior to the local model fine-tuning phase. The central concept is to perturb local data by injecting Gaussian-distributed noise to obfuscate sensitive features, adhering to differential privacy at the level of participants while preserving acceptable learning utility.

After the federated Nash Bargaining has completed (providing fair contribution-weighted allocations to each participant in the model update), each local client has a portion of the model update budget available to them. At this point, a differential privacy budget, ϵ is available for each client to satisfy the sensitivity and privacy constraints for their respective data environment.

Let:

- D_i be the original dataset of client i
- \tilde{D}_i be the noise-perturbed dataset
- $\mathcal{N}(0, \sigma^2)$ be the Gaussian distribution with mean 0 and variance σ^2
- ϵ be the differential privacy parameter
- Δf be the sensitivity of the data (maximum change in output caused by a single data point change)

The Gaussian noise added to each data point is given by:

$$\tilde{D}_i = D_i + \mathcal{N}(0, \sigma^2)$$

Where σ is calibrated using:

$$\sigma \geq \frac{\Delta f \sqrt{2 \log(1.25/\delta)}}{\epsilon}$$

The noise is independently sampled and added to each feature vector $\mathbf{x}_{ij} \in D_i$, for each data point j :

$$\tilde{\mathbf{x}}_{ij} = \mathbf{x}_{ij} + \mathbf{z}_{ij}, \quad \text{where } \mathbf{z}_{ij} \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$$

This perturbation strategy protects input data from reverse engineering because it creates obstacles that stop attackers from reconstructing the original data during resource-limited or sensitive IIoT applications as well as for other Critical Infrastructures.

This step plays a crucial role in:

- Limiting the exposure of raw feature patterns during local computations
- Ensuring compliance with **local differential privacy (LDP)** constraints
- Enhancing robustness of the federated IDS against inversion or reconstruction attacks

Once local training is completed using \tilde{D}_i , the model update is clipped and encrypted before being pushed back to the central server for aggregation.

3.5 Robust Aggregation

3.5.1 Adversarial Attacks

In federated learning (FL) environments, adversarial clients pose a unique challenge to convergence and system integrity. In particular, adversarial clients may produce faulty updates accidentally, or intentionally as part of a coordinated attack, corrupting the global model by producing hostile updates that are included in the aggregation. Of all possible adversarial behavior patterns, Byzantine attacks are considered one of the most disruptive.

3.5.1.1 Byzantine Attacks in Federated Learning

Byzantine participants are the federated network participants that submit arbitrary updates, and deviates from the expected protocol, and/or submit strategically poisoned updates. Byzantine updates may consist of:

- Random noise to composite the global model (e.g., Random attacks)
- Sign-flipping to reverse the gradient direction
- Scaling attacks to exaggerate local gradients
- Targeted poisoning to misclassify inputs (e.g., backdoor attacks)

The term Byzantine comes from the Byzantine Generals Problem. In the Byzantine Generals Problem, some participants may act antagonistically or inexplicably leading to breakdown of consensus in a decentralized system.

In a typical FL setup, the global server aggregates updates from N clients as:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \eta \cdot \frac{1}{N} \sum_{i=1}^N \Delta \mathbf{w}_i$$

Under a Byzantine scenario, if a subset $B \subset N$ of clients are adversarial, their updates $\Delta \mathbf{w}_i$ may be arbitrarily corrupted. As a result, the aggregated model \mathbf{w}_{t+1} can diverge or perform sub-optimally.

3.5.2 Byzantine Fault Tolerant Aggregation

To mitigate the vulnerabilities introduced by adversarial participants in federated learning, namely Byzantine behavior, Byzantine Fault Tolerant (BFT) aggregation techniques have emerged. BFT aggregation techniques help to ensure that the global model can converge robustly under the influence of some fraction of malicious or unreliable clients.

Problem Definition:

Let $S = \{\Delta \mathbf{w}_1, \Delta \mathbf{w}_2, \dots, \Delta \mathbf{w}_N\}$ denote the set of local model updates received by the server from N clients. Assume up to $f < N$ of these updates may be arbitrarily malicious. The goal of BFT aggregation is to compute an aggregated update $\Delta \mathbf{w}_{agg}$ that minimizes the influence of these faulty clients.

In the proposed framework, we incorporate **Multi-Krum** as a defense mechanism during global model updates with integrated Reputation-based weighing. This helps mitigate the effects of adversarial updates, ensure model integrity, and support stable convergence even in the presence of partial client corruption.

3.5.2.1 Multi-Krum Aggregation

Multi-Krum is a robust aggregation algorithm designed to mitigate Byzantine failures in Federated Learning. It selects a subset of client updates based on Euclidean distance metrics to ensure secure aggregation.

→ **Inputs**

- **updates:** A set of update vectors from multiple clients, denoted as $U = \{u_1, u_2, \dots, u_n\}$
- **num_selected:** The number of client updates to be selected for aggregation, denoted as k .

→ **Compute Pairwise Euclidean Distances**

For each client update u_i , compute the Euclidean distance to every other client update u_j :

$$d_{ij} = \|u_i - u_j\|_2 = \sqrt{\sum_{m=1}^M (u_{i,m} - u_{j,m})^2}, \quad \forall i \neq j$$

where M is the dimension of the update vector.

→ **Compute Selection Scores**

For each client i , sort the distances d_{ij} in ascending order and select the k nearest neighbors:

$$S_i = \sum_{j \in N_k(i)} d_{ij}$$

where $N_k^{(i)}$ represents the indices of the k nearest clients to i .

→ **Select the Best Clients**

The clients with the lowest selection scores are chosen for aggregation:

$$C = \operatorname{argmin}_i S_i, \quad \text{for } i \in \{1, \dots, n\}$$

→ **Compute Aggregated Update**

The final aggregated update is computed as the mean of the selected updates:

$$u_{agg} = \frac{1}{k} \sum_{i \in C} u_i$$

3.5.2.2 Reputation-based Weighing

To improve robustness against malicious and unreliable clients in federated learning, we introduce a reputation-based weighting means in conjunction with Multi-Krum aggregation. While, on the one hand, Multi-Krum can reject outliers and adversarial updates using distance-based scoring, it does not account for historical trust or consistency with the client behavior over training rounds. This is where the reputation-based approach can become very useful.

Multi-Krum relies only on the distance amongst model updates and can eliminate a few Byzantine (malicious or faulty) updates, but never learns which clients are reliably and unreliably faulty. This means that Multi-Krum can be easily used against adaptive adversaries, who will manipulate their updates to stay within the distance limitation. We use reputation-based aggregation as a second layer of filtering:

- ✓ We use reputation-based aggregation as a second layer of filtering: Multi-Krum allows for an initial selection to identify plausible updates.
- ✓ The reputation system adds temporal memory, rewarding honest participation and acts to collectively punish anomalous behavior over time.
- ✓ Both techniques are stronger together in that they provide spatial and temporal robustness and can help deter stealthy or slow poisoning attacks.

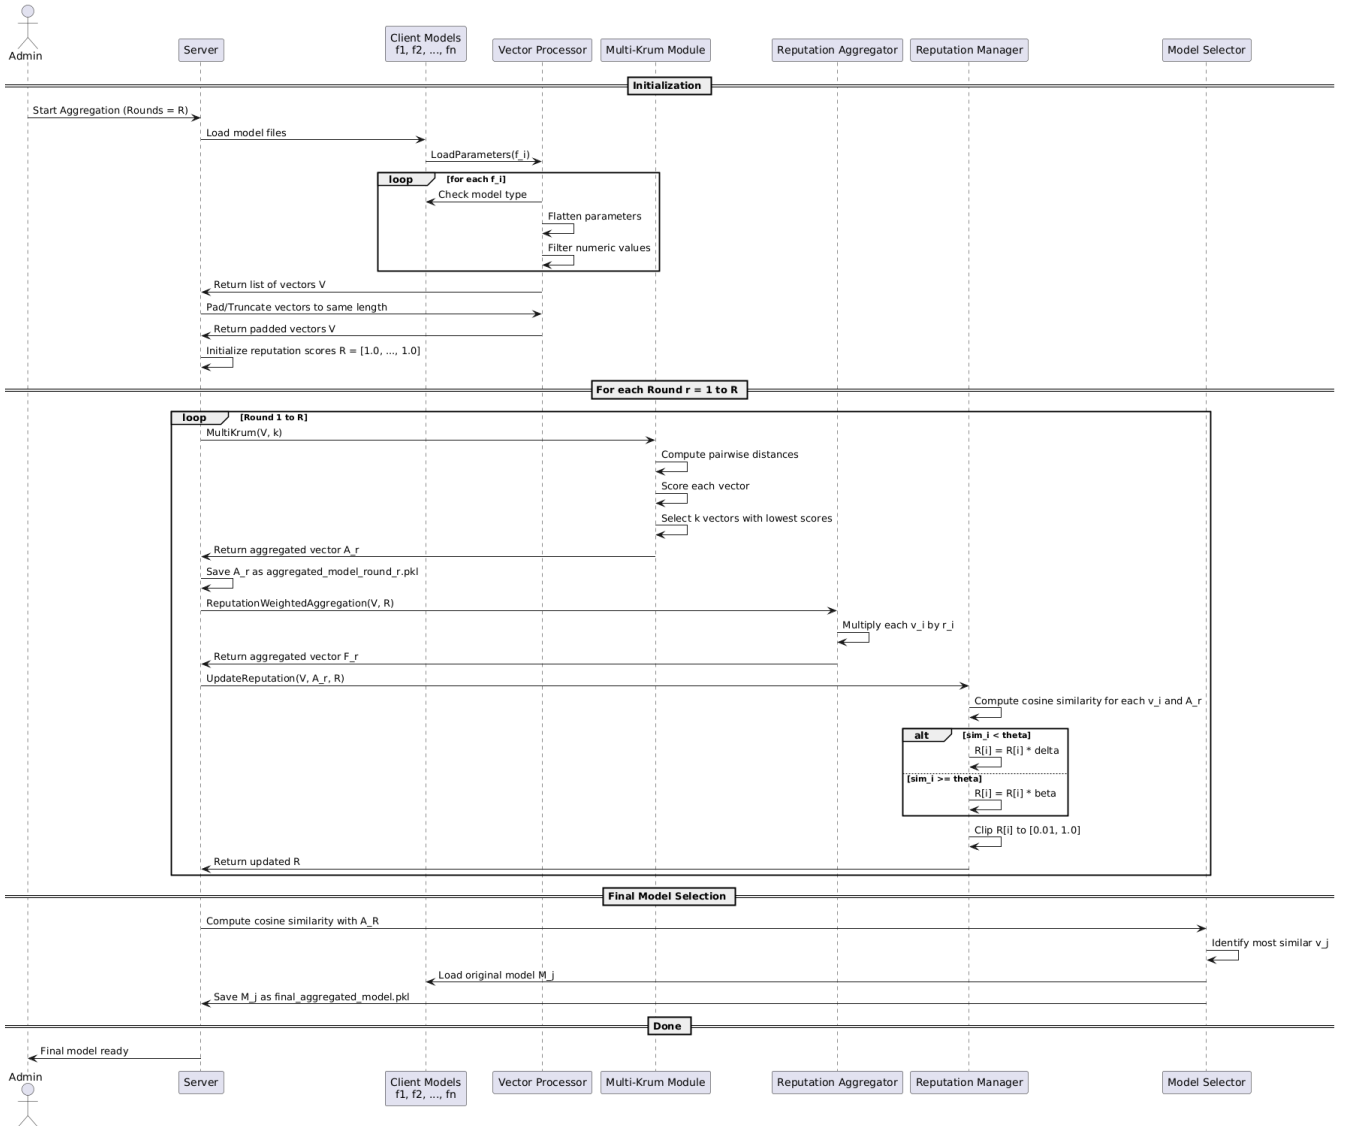


Figure 3.2 BFT Aggregation using Multi-Krum and Reputation Weighting

There is also the additional benefit of a hybrid, reliable and adaptable approach to Byzantine Fault Tolerance in federated learning for generally increasing the security of model convergence in more adverse environments [Figure 3.2].

3.6 Iterative Model Updates

Clients receive the latest global model post-aggregation, iteratively enhancing the system's capability to detect evolving intrusion patterns in real time [Figure 3.3].

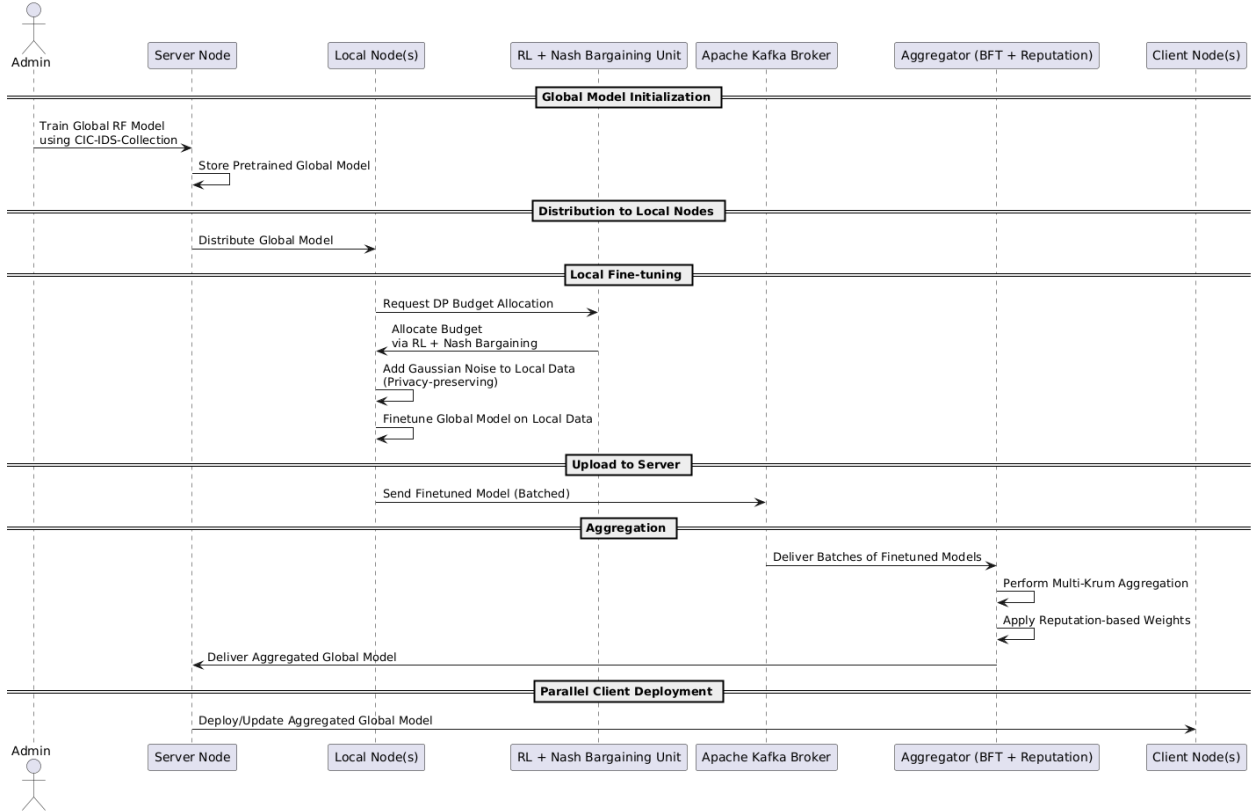


Figure 3.3 High-Level Sequence of the Proposed Methodology

3.7 Quantum-Resilient Communication

To ensure the security and confidentiality of federated model parameters during transmission—especially in the face of emerging quantum threats—we employ a hybrid encryption mechanism that combines **Advanced Encryption Standard (AES-256)** for data confidentiality with simulated **Kyber-based key encapsulation** for quantum-resilient session key exchange. This two-tier cryptographic mechanism mimics post-quantum key exchange while using symmetric encryption for actual payload security.

3.7.1 Symmetric Encryption using AES-256

AES-256 in Cipher Block Chaining (CBC) mode is used to encrypt serialized model files (e.g., .pkl files). This involves:

- Generation of a 256-bit (32-byte) **random AES key** K_{AES} .
- Generation of a **random Initialization Vector (IV)** of 128 bits (16 bytes), denoted as IV .

The model file M is first padded using **PKCS#7-style padding** (manual space padding in our simulation) and then encrypted using:

$$C = \text{AES}_{\text{CBC}}(K_{\text{AES}}, IV, \text{Pad}(M))$$

Where:

→ C = Ciphertext

→ $\text{Pad}(M)$ = Padded binary data of the model file

The encryption function:

```
cipher = AES.new(aes_key, AES.MODE_CBC, iv)
encrypted_data = cipher.encrypt(padded_data)
```

3.7.2 Simulated Kyber Key Encapsulation

To simulate a **quantum-resistant key exchange**, we implement a mock version of **Kyber Key Encapsulation Mechanism (KEM)**, aligned with NIST PQC standardization.

Key Generation:

→ A simulated public key $PK \in \{0,1\}^{256}$

→ A simulated secret key $SK \in \{0,1\}^{512}$

$$(PK, SK) = \text{Kyber.KeyGen}()$$

Key Encapsulation:

A symmetric session key $K_{\text{session}} \in \{0,1\}^{256}$ is generated and encapsulated into an encrypted form $C_{\text{kyber}} \in \{0,1\}^{512}$ using the simulated public key:

$$(K_{\text{session}}, C_{\text{kyber}}) = \text{Kyber.Encapsulate}(PK)$$

This session key can later be recovered only by the node possessing the corresponding secret key.

3.7.3 Secure Transmission Structure

The encrypted data package contains:

- C (encrypted_data)
- IV used for AES
- C_{kyber} (encapsulated_key)
- **Public key** (for validation or audit trails)
- K_{session} (session key that is generated during the Kyber key encapsulation process)

3.7.4 Decryption and De-Encapsulation

Upon receiving the encrypted payload, the federated server or edge node:

- Decrypts the AES-encrypted data using the K_{session} , recovered from the encapsulated key using:

$$K_{\text{session}} = \text{Kyber.Decapsulate}(SK, C_{\text{kyber}})$$

- Decrypts the ciphertext using AES-CBC:

$$M' = \text{AES}_{\text{CBC}}^{-1}(K_{\text{session}}, IV, C)$$

- Removes padding to obtain original model file M .

3.7.5 Security Analysis and Quantum Resilience

- AES-256 is currently considered secure against quantum attacks due to the large key size, where Grover's algorithm only yields quadratic speedup, effectively reducing the security level to 128-bit—still considered strong.
- Kyber (Lattice-based encryption) is resistant to Shor's algorithm and is one of the selected algorithms in NIST's post-quantum cryptography standardization.

→ This hybrid encryption setup ensures forward secrecy, integrity, and confidentiality under both classical and quantum threat models.

Equation Summary

Let:

- $K_{AES} \leftarrow \text{RBG}(256)$
- $IV \leftarrow \text{RBG}(128)$
- $(PK, SK) \leftarrow \text{Kyber.KeyGen}()$
- $(K_{session}, C_{kyber}) \leftarrow \text{Kyber.Encapsulate}(PK)$

Then:

$$C = \text{AES}_{CBC}(K_{AES}, IV, \text{Pad}(M))$$

$$\text{Encrypted_Object} = \{IV, C, PK, C_{kyber}\}$$

Decryption:

$$K_{session} = \text{Kyber.Decapsulate}(SK, C_{kyber})$$

$$M = \text{Unpad}(\text{AES}_{CBC}^{-1}(K_{session}, IV, C))$$

Chapter-4

DEPLOYMENT

4.1 Front-End Deployment

4.1.1 Core Libraries & Frameworks

- **React:**
 - Component-based architecture for reusable UI elements.
 - Efficient state updates via virtual DOM.
- **React Router (react-router-dom):**
 - Declarative routing for seamless navigation (e.g., Dashboard, Network Map).
 - Dynamic page rendering without full reloads.
- **Redux Toolkit:**
 - Manages global state (e.g., flow tracking, alerts).
 - Reduces boilerplate for complex state logic.
- **@tanstack/react-query:**
 - Handles server-state (API calls, caching, re-fetching).
 - Abstracts loading/error states for components.
- **Webhooks:**
 - Streams real-time flow data and alerts.
 - Enables live monitoring without delays.

4.1.1.1 Styling & Visuals

- **Tailwind CSS:**
 - Utility-first framework for rapid styling (e.g., flex, p-4).
 - Ensures responsive, modern UI.
 - **Recharts:**
 - Interactive charts for data visualization (e.g., traffic volumes, protocols).
 - **react-icons:**
 - Provides recognizable icons (e.g., Font Awesome pack).
 - **A-Frame & aframe-extras:**
 - 3D visualization for network topologies.
 - **react-force-graph:**
 - Physics-based network graphs for device connections.
-

4.1.2 Development & Tooling

- **Vite:**
 - Fast build tool with HMR for developer productivity.
-

4.1.3 Application Workflow

4.1.3.1 Initialization

- **Main.jsx mounts <App /> with providers:**
 - Redux Provider for global state.
 - React Query QueryClientProvider for server-state.
 - React Router BrowserRouter for navigation.
- **Webhooks connection initialized for real-time data.**

4.1.3.2 Page Navigation (e.g., /dashboard)

- React Router renders Dashboard.
- **Actions:**
 - Fetches historical data via API (fetchRecentFlows).
 - Accesses live data from Redux (flowsSlice, alertsSlice).

4.1.3.3 Real-Time Updates

- **WebhooksService:**
 - Enriches raw flows (e.g., computes total_bytes).
 - Dispatches addFlow to Redux.
 - Triggers addAlert and toast notifications for anomalies.

4.1.3.4 User Interactions

- View dashboards, filter alerts, investigate flows.
- Updates Redux (global) or component state (local).

4.1.4 Dataflow

4.1.4.1 Inputs

- Webhooks: Live flow data/alerts.
- REST API: Historical data.
- User Interactions: Clicks, filters, searches.

4.1.4.2 Transformations

- WebhooksService: Computes fields (e.g., total_bytes).
- Redux: Filters/sorts data.
- UI: Converts raw JSON to visualizations.

4.1.4.3 Data Storage

- **Redux:**
 - flowsSlice: Live/historical flows.
 - alertsSlice: Active alerts.
- **React Query Cache: API responses (e.g., fetchRecentFlows).**

4.1.4.4 Outputs

- **Charts, tables, toast notifications, network maps.**

4.1.5 Common Use Cases

- Real-Time Monitoring: Spot unusual activity (e.g., suspicious IPs).
- Alert Management: Filter/dismiss alerts by severity.
- Flow Investigation: View IPs, bytes, protocols.
- Historical Analysis: Identify traffic patterns.

4.1.6 Deployment

- Platform: AWS/Azure/Vercel.
 - Static Assets: Vite + CDN.
 - CI/CD: GitHub Actions/CircleCI.
 - Security: WSS, input sanitization, authentication.
-

4.1.7 Testing & Observability

- Unit Tests: Redux reducers (Jest).
- Integration Tests: Components + Redux (React Testing Library).
- E2E Tests: User journeys (Cypress).
- Monitoring: Redux DevTools, Sentry (optional).

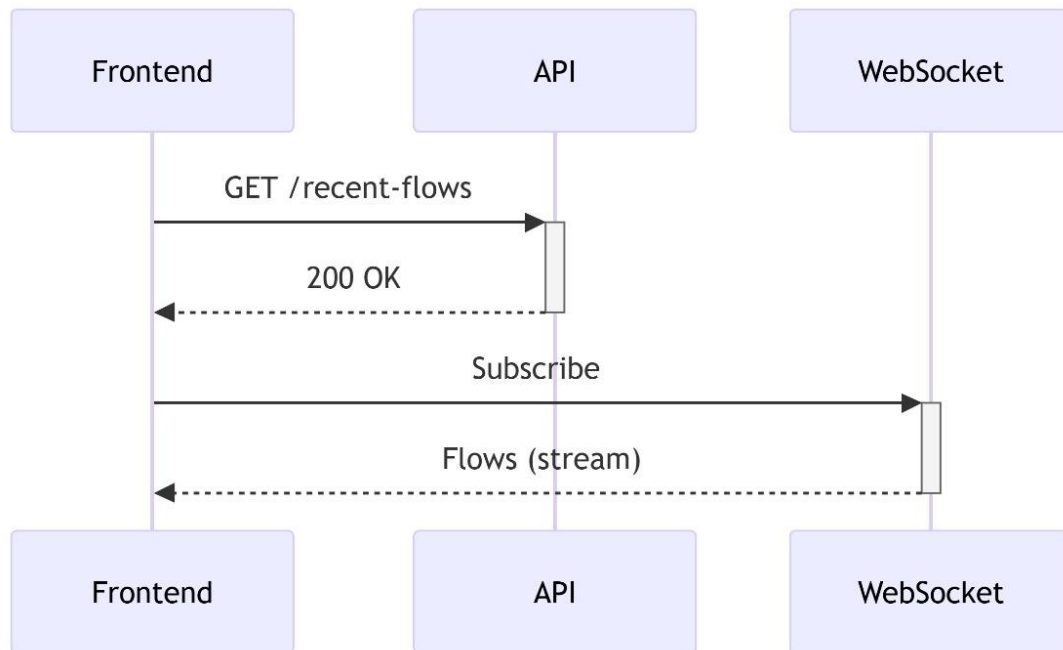


Figure 4.1 Front-End Workflow

4.2 Back-End Deployment

4.2.1 Core Infrastructure

- **Zookeeper:**
 - Manages Kafka cluster coordination and maintains configuration data.
 - Configured via `zookeeper.properties` (port 2182, data directory).
- **Kafka:**
 - Acts as the central message broker for streaming network flow data.
 - Configured via `server.properties` (broker ID, listeners, log directories, Zookeeper connection).
 - Uses topic `network-flows` to distribute flow data to consumers.

4.2.2 Data Producer

Network Sniffer (`kafka_produce_capture.py`):

- **Role:** Captures network traffic (via Scapy), aggregates packets into flows, and publishes flow metrics to Kafka.

- **Key Features:**
 - Flow tracking (bidirectional, TCP/UDP/ICMP).
 - Kafka integration with retries and batching.
 - Metrics include packet counts, byte sizes, TCP flags, and flow duration.
- **Input:** Network interface (e.g., wlo1).

Output: Kafka topic *network-flows*.

4.2.3 Data Consumers

4.2.3.1 Redis TimeSeries Consumer (*kafka_consumer_redis.py*):

- Role: Subscribes to Kafka, ingests flow data into Redis TimeSeries for real-time analytics.
- Key Features:
 - Stores metrics like total bytes, packets, TCP flags, and protocol-specific stats.
 - Downsampling and retention policies for time-series data (1s, 10s, 1m resolutions).
- Storage: Redis with TimeSeries module (via Docker).

4.2.3.2 PostgreSQL Consumer (*kafka_consume_db.py*):

- Role: Persists flow data into PostgreSQL for structured querying and long-term storage.
- Key Features:
 - Two tables: raw JSON (*network_flows*) and structured metrics (*flow_metrics*).
 - Indexes for fast querying on IPs, ports, and timestamps.
- Schema: Includes fields like *src_ip*, *dst_ip*, *total_bytes*, and *timestamp*.

4.2.3.3 FastAPI Consumer (*kafka_consume_api.py*):

- Role: Provides a real-time API to access the latest flow data.
- Key Features:
 - In-memory storage of recent flows (last 1,000 entries).
 - Endpoints: */flows*, */flows/latest*, */flows/search*.
 - Enriches data with *total_bytes* for API responses.

Runtime: ASGI server (Uvicorn) on port 8888.

4.2.4 Supporting Components

- Docker Compose:
 - Orchestrates Zookeeper and Kafka containers.
 - Maps volumes for data persistence (*./kafka-data*).
- PostgreSQL/Redis:
 - External services (configured via CLI arguments).
 - PostgreSQL credentials: *network_admin/securepassword123*.
 - Redis runs via *redislabs/redistimeseries* Docker image.

4.2.5 Data Flow

- Producer → Kafka: Network flows are captured and sent to Kafka.
- Kafka → Consumers: All three consumers read from the same Kafka topic in parallel.
- Consumers → Storage/API:
 - Redis: Optimized for time-series queries (e.g., traffic spikes).
 - PostgreSQL: Structured analytics (e.g., historical trends).
 - FastAPI: Real-time monitoring via HTTP.

4.2.6 Technologies Used

- Messaging: Apache Kafka (with Zookeeper).
- Data Capture: Scapy (packet sniffing).
- Storage: Redis TimeSeries, PostgreSQL.
- API: FastAPI (Python), Uvicorn.
- Orchestration: Docker Compose.

This architecture enables scalable network monitoring with real-time analytics, historical storage, and API access. Each component is decoupled, allowing independent scaling (e.g., adding more consumers or partitioning Kafka topics) [Table 4.1].

Table 4.1 Technologies

Technology	Role
Apache Kafka	Distributed message broker for streaming network flow data between producers and consumers.
Apache Zookeeper	Manages Kafka cluster coordination, leader election, and configuration.
PostgreSQL	Relational database for structured storage of flow metrics (e.g., IPs, ports, protocols).
Redis (TimeSeries Module)	Time-series database for real-time analytics (e.g., traffic spikes, protocol trends).
ASGI (Uvicorn)	Asynchronous server protocol to run the FastAPI application efficiently.

Programming language used in backend is python with various libraries and frameworks [Table 4.2].

Table 4.2 Libraries and frameworks

Library/Framework	Role
Scapy	Captures and analyzes network packets; extracts flow metrics (TCP/UDP/ICMP).
Kafka-Python	Python client for producing/consuming Kafka messages.
psycopg2-binary	PostgreSQL adapter for Python; inserts flow data into structured tables.
Redis/redis-py-cluster	Python client for interacting with Redis TimeSeries.
FastAPI	Web framework for building real-time API endpoints to query flow data.
Pydantic	Validates API request/response data models in FastAPI.
NumPy	Computes statistical metrics (e.g., mean packet lengths) in flow data.
argparse	Parses command-line arguments for script configuration.
logging	Standard Python module for application logging and debugging.

Table 4.3 Tools & Infrastructure

Tool	Role
Docker	Containerizes Kafka, Zookeeper, and Redis for isolated deployment.
Docker Compose	Orchestrates multi-container setup (Kafka + Zookeeper).
Kafka CLI Tools	Manages Kafka topics (e.g., kafka-topics --create).
PostgreSQL CLI (psql)	Interacts with the PostgreSQL database via terminal.
Uvicorn	ASGI server to run the FastAPI application.
Gunicorn	Optional production-grade server for FastAPI (mentioned in requirements).
venv	Creates Python virtual environments for dependency isolation.

Table 4.4 Data Formats & Protocols

Format/Protocol	Role
JSON	Serializes flow data for Kafka messages and PostgreSQL JSONB storage.
HTTP	Protocol for FastAPI endpoints to serve flow data.
TCP/UDP/ICMP	Network protocols analyzed by Scapy during packet capture.

Key Interactions

- **Scapy** → **Kafka-Python**: Captured packets are aggregated into flows and sent to Kafka.
- **Kafka-Python** → **PostgreSQL/Redis**: Consumers ingest Kafka messages into databases.
- **FastAPI** → **Redis/PostgreSQL**: API queries databases to serve real-time flow data.
- **Docker Compose** → **Kafka/Zookeeper**: Simplifies deployment of messaging infrastructure.

This stack enables scalable network monitoring with real-time analytics, historical storage, and API access. Each component is modular, allowing independent scaling (e.g., adding more Kafka partitions or Redis instances).

4.2.7 End-to-End Workflow of the Application [Figure 4.1]

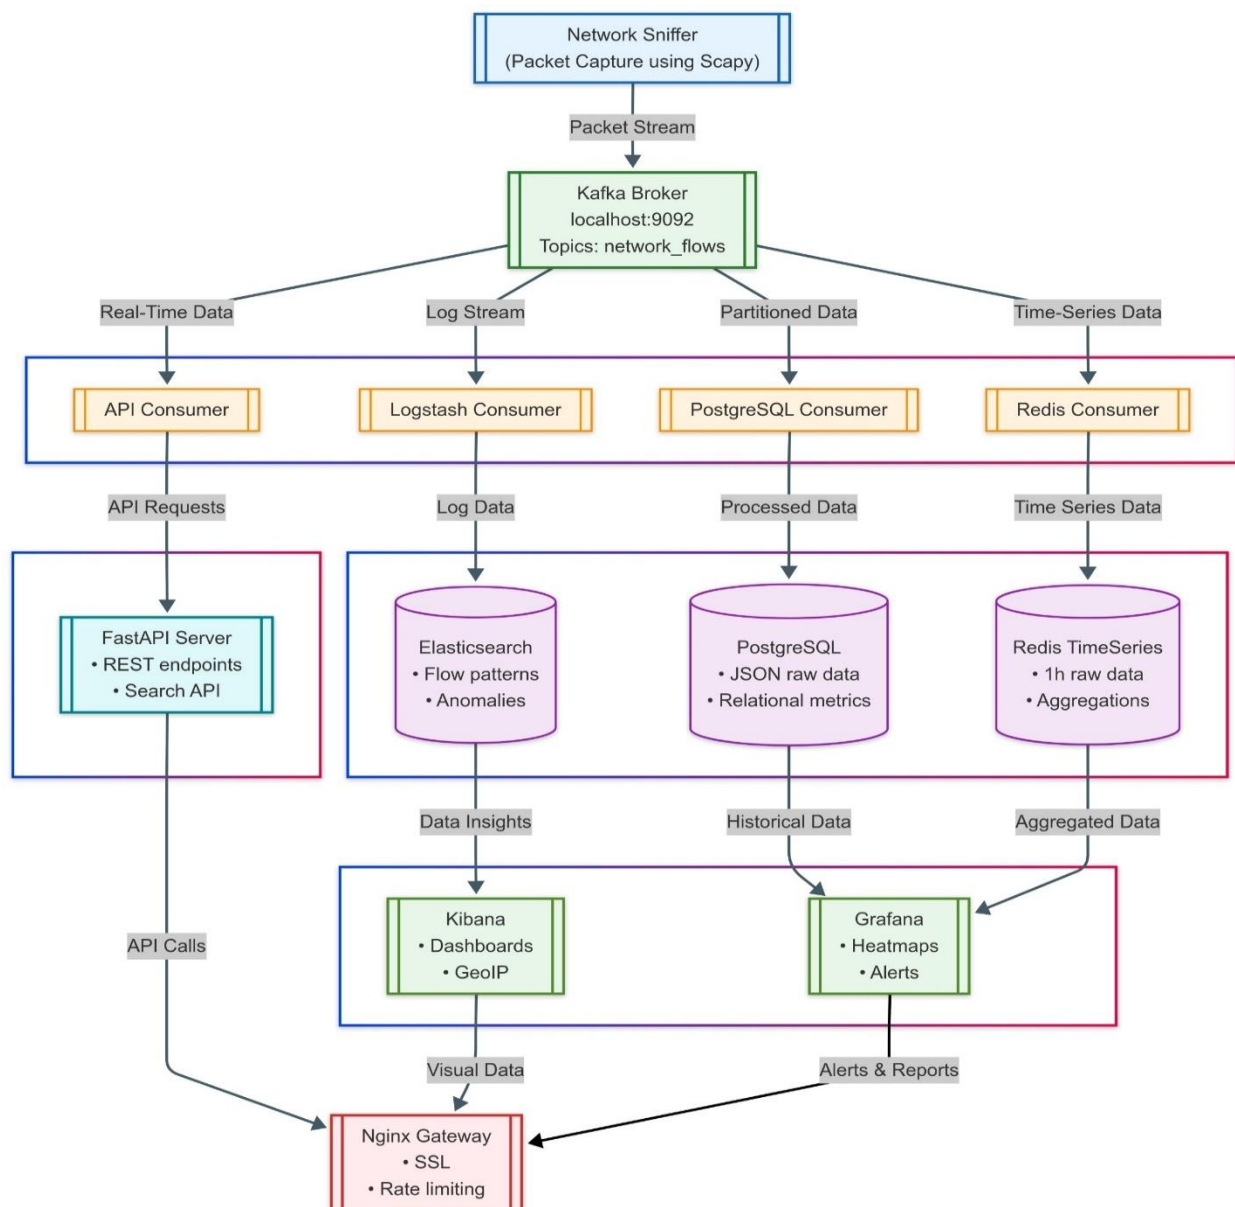


Figure 4.2 End-to-End Workflow

1. Infrastructure Setup

→ Start Zookeeper & Kafka:

- Use docker-compose.yml to launch Zookeeper and Kafka containers.
- Kafka brokers are configured to use Zookeeper for cluster coordination.
- Kafka topic network-flows is created with 3 partitions and replication factor 1.

→ Initialize Databases:

- PostgreSQL: Create netflows database and tables via kafka_consume_db.py.
- Redis: Start Redis TimeSeries Docker container (redislabs/redistimeseries).

2. Data Production

→ Packet Capture:

- Run kafka_produce_capture.py with a network interface (e.g., wlo1).
- Scapy captures live network traffic and groups packets into **bidirectional flows** (based on IPs, ports, protocol).

→ Flow Aggregation:

- Track metrics per flow: duration, packets, bytes, TCP flags (SYN/FIN/RST), etc.
- Flows expire after 120s of inactivity or explicit termination (e.g., TCP RST/FIN).

→ Publish to Kafka:

- Serialize flow metrics into JSON and send to Kafka topic network-flows.
- Kafka producer batches messages for efficiency and handles retries on errors.

3. Data Consumption

3.1 PostgreSQL Consumer

→ **Role:** Persist structured flow data for long-term analytics.

→ Process:

- Subscribe to Kafka topic network-flows.
- For each message:
 - i. Insert raw JSON into network_flows table.
 - ii. Parse and store structured metrics (IPs, ports, protocols) in flow_metrics table.
 - iii. Create indexes for fast querying (e.g., src_ip, timestamp).

3.2 Redis TimeSeries Consumer

→ **Role:** Enable real-time time-series analytics (e.g., traffic spikes).

→ **Process:**

- Subscribe to Kafka topic network-flows.
- For each flow:
 - i. Store metrics (bytes, packets, TCP flags) in Redis TimeSeries with labels (IP, protocol).
 - ii. Downsample data for retention (1s → 10s → 1m resolutions).
 - iii. Create compaction rules for aggregation (e.g., average bytes per minute).

3.3 FastAPI Consumer

→ **Role:** Serve real-time flow data via HTTP API.

→ **Process:**

- Subscribe to Kafka topic network-flows (latest offset).
- Maintain in-memory buffer (deque) of the last 1,000 flows.
- Expose endpoints:
 - i. /flows: List recent flows (paginated).
 - ii. /flows/latest: Latest flow entry.
 - iii. /flows/search: Filter by IP/protocol.

4. Data Querying

→ **PostgreSQL:**

- Use SQL queries to analyze historical data (e.g., "Top 10 source IPs by bytes").

```
SELECT src_ip, SUM(total_bytes)
FROM flow_metrics
GROUP BY src_ip
ORDER BY SUM(total_bytes) DESC
LIMIT 10;
```

→ **Redis TimeSeries:**

- Query time-series metrics (e.g., "Bytes per protocol in the last hour"):

```
TS.RANGE ts:protocol_bytes - + AGGREGATION avg 3600000
```

→ **FastAPI:**

- Access real-time data via HTTP:

`curl http://localhost:8888/flows/search?src_ip=192.168.1.10`

5. Monitoring & Scaling

- **Kafka Monitoring:**

- Use `kafka-console-consumer.sh` to debug message flow:

`kafka-console-consumer --bootstrap-server localhost:9092 --topic network-flows`

- **Scaling:**

- Increase Kafka partitions to parallelize consumer workloads.
- Deploy consumers as microservices (e.g., Kubernetes pods).

1. Zookeeper & Kafka (Docker Compose)

Purpose:

- **Zookeeper:** Manages Kafka cluster coordination (broker election, topic configuration).
- **Kafka:** Acts as the central message bus for streaming network flow data.

Internal Logic:

- Defined in `docker-compose.yml`:
 - Zookeeper runs on port 2182 with persisted data in `./kafka-data/zookeeper`.
 - Kafka broker connects to Zookeeper and exposes port 9092.
 - Topics are created manually (e.g., `network-flows`).

Integration:

- All producers/consumers connect to Kafka via `localhost:9092`.
- Kafka uses Zookeeper for cluster health checks and metadata storage.

2. Network Sniffer Producer (`kafka_produce_capture.py`)

Purpose: Capture network packets, aggregate flows, and publish to Kafka.

Internal Logic:

→ **Packet Capture:**

- Uses `scapy` to sniff packets on a network interface (e.g., `wlo1`).
- Classifies packets into **bidirectional flows** (sorted by IP/port/protocol).

→ **Flow Tracking:**

- Maintains a dictionary of active flows (self.flows).
- Tracks metrics: packets, bytes, TCP flags, timestamps.
- Flows expire after 120s or on TCP FIN/RST.

→ **Kafka Publishing:**

- Serializes flow stats to JSON.
- Uses KafkaProducer with batching (linger_ms=100) and retries.

Integration:

- Sends data to Kafka topic network-flows.
- Downstream consumers (PostgreSQL, Redis, FastAPI) subscribe to this topic.

3. PostgreSQL Consumer (kafka_consume_db.py)

Purpose: Persist flow data for structured analytics.

Internal Logic:

→ **Kafka Consumption:**

- Uses KafkaConsumer in group_id='postgres-consumer-group'.
- Processes messages in batches with poll(timeout_ms=1000).

→ **Database Operations:**

- **Raw Storage:** Inserts JSON flow data into network_flows (JSONB column).
- **Structured Storage:** Parses JSON into flow_metrics (columns: src_ip, total_bytes, etc.).
- **Indexes:** Created on src_ip, dst_ip, timestamp for fast querying.

→ **Retry Logic:**

- Reconnects to PostgreSQL on failure (max 5 retries).

Integration:

- Connects to PostgreSQL using psycopg2 (credentials from CLI args).
- Enables SQL-based analytics (e.g., "Top talkers by traffic volume").

4. Redis TimeSeries Consumer (kafka_consumer_redis.py)

Purpose: Enable real-time time-series analytics.

Internal Logic:**→ Kafka Consumption:**

- Subscribes to network-flows in `group_id='redis-timeseries-consumer-group'`.

→ TimeSeries Storage:

- Uses Redis commands (TS.ADD, TS.CREATE, TS.CREATERULE).
- Stores metrics:
 - flow_bytes_total: Total bytes per flow (labels: src_ip, protocol).
 - tcp_syn_count: SYN flags per source IP.
- Downsampling rules aggregate data into 10-second and 1-minute buckets.

→ Pipeline Optimization:

- Batches Redis operations for efficiency.

Integration:

- Connects to Redis via `redis-py`.
- Enables queries like "Bytes per protocol in the last 5 minutes" via Redis CLI.

5. FastAPI Consumer (kafka_consume_api.py)

Purpose: Provide real-time API access to flow data.

Internal Logic:**→ Kafka Consumption:**

- Runs in a background thread (`start_kafka_consumer`).
- Maintains a rotating buffer (`deque(maxlen=1000)`) of recent flows.

→ API Endpoints:

- `/flows`: Returns last N flows (paginated).
- `/flows/search`: Filters flows by IP/protocol.
- Enriches data with `total_bytes = fwd_bytes + bwd_bytes`.

→ Concurrency:

- Uses ASGI server (`uvicorn`) for asynchronous request handling.

Integration:

- Subscribes to Kafka and serves HTTP requests on port 8888.
- Integrates with frontend dashboards for real-time monitoring.

6. Supporting Components**Docker Compose**

- **Role:** Deploy Kafka/Zookeeper with persistent storage.
- **Key Configs:**
 - Volume mounts for Kafka/Zookeeper data.
 - Environment variables for broker ID, listeners, and replication.

PostgreSQL Schema

- **Tables:**
 - network_flows: Raw JSON data.
 - flow_metrics: Structured columns (IPs, ports, bytes).
- **Indexes:** Optimized for common queries (e.g., src_ip).

Redis TimeSeries

- **Data Retention:**
 - Raw data: 1 hour (1-second resolution).
 - Aggregated data: 1 day (10-second), 1 week (1-minute).

Component Interactions**→ Producer → Kafka:**

- Network flows are serialized to JSON and published to network-flows.

→ Kafka → Consumers:

- PostgreSQL/Redis/FastAPI consumers read from the same topic in parallel.
- Each consumer group processes messages independently.

→ Data Storage → API:

- FastAPI serves data from its in-memory buffer (not directly from DBs).
- PostgreSQL/Redis are queried separately for historical/analytical use cases.

Error Handling & Scaling

- **Kafka:** Retries on producer/consumer errors.
- **PostgreSQL/Redis:** Reconnection logic for database failures.
- **Scaling:**
 - Kafka: Add partitions to distribute load.
 - Consumers: Deploy multiple instances (e.g., Kubernetes pods).

This modular architecture ensures fault tolerance, real-time processing, and scalability. Each component can be upgraded or replaced independently (e.g., swapping Scapy for libpcap-based capture).

Data Flow with ELK Integration:

→ **Producer** → **Kafka** (Unchanged):

- The `kafka_produce_capture.py` script captures network flows, aggregates metrics, and publishes JSON data to the Kafka topic `network-flows`.

→ **Kafka** → **Logstash**:

- Logstash uses its **Kafka input plugin** to consume messages from `network-flows`.

Example Logstash pipeline:

```
input {
  kafka {
    bootstrap_servers => "localhost:9092"
    topics => ["network-flows"]
    codec => json
  }
}
```

→ **Logstash** → **Elasticsearch**:

- Logstash processes data (e.g., timestamp conversion, geo-IP enrichment) and sends it to Elasticsearch:

```
filter {
  date { match => [ "timestamp", "UNIX" ] } # Convert UNIX timestamp to
  ISO format
```

```
# Optional: Add geo-IP data
geoip { source => "src_ip" target => "geoip" }
}

output {
  elasticsearch {
    hosts => ["http://elasticsearch:9200"]
    index => "network-flows-%{+YYYY.MM.dd}"
  }
}
```

→ **Elasticsearch ↔ Kibana:**

- Kibana connects to Elasticsearch to create dashboards for:
 - i. Real-time traffic volume
 - ii. Top source/destination IPs and ports
 - iii. Protocol distribution (TCP/UDP/ICMP)
 - iv. Geo-IP heatmaps (if geo-IP filter is enabled)
 - v. Alerting on traffic spikes or anomalies

Key Inputs, Transformations, and Outputs:

Component	Role
Input	JSON flow data from Kafka (network-flows topic).
Transformations	<ul style="list-style-type: none"> - Timestamp normalization - Geo-IP enrichment (optional) - Field type mapping (e.g., protocol to keyword).
Data Stores	<ul style="list-style-type: none"> - Elasticsearch: Indexed flow data for fast search/analytics - PostgreSQL/Redis: Existing structured/time-series data.
Output	<ul style="list-style-type: none"> - Kibana: Interactive dashboards - Grafana (optional): Integrate Elasticsearch as a data source.

Component Interactions:

- **Kafka:** Serves as the central hub for distributing flow data to all consumers (PostgreSQL, Redis, FastAPI, Logstash).
- **Logstash:** Processes data in parallel with other consumers, ensuring no impact on existing workflows.
- **Elasticsearch:** Complements PostgreSQL (structured analytics) and Redis (real-time metrics) by enabling full-text search, complex aggregations, and long-term log retention.
- **Kibana:** Provides a user-friendly interface for security analysts to explore network traffic patterns.

Deployment Considerations:

→ Docker Integration:

- Extend docker-compose.yml to include Elasticsearch, Logstash, and Kibana containers:

services:

elasticsearch:

image: docker.elastic.co/elasticsearch/elasticsearch:8.12.0

ports: ["9200:9200"]

environment:

– discovery.type=single-node

logstash:

image: docker.elastic.co/logstash/logstash:8.12.0

volumes:

– ./logstash/pipeline:/usr/share/logstash/pipeline

kibana:

image: docker.elastic.co/kibana/kibana:8.12.0

ports: ["5601:5601"]

→ **Data Retention:**

- Use Elasticsearch Index Lifecycle Management (ILM) to automatically roll over and delete old indices (e.g., retain data for 30 days).

→ **Security:**

- Enable TLS/SSL for Elasticsearch and authentication for Kibana.

Benefits:

- **Enhanced Analytics:** Kibana dashboards enable ad-hoc queries and anomaly detection.
- **Scalability:** Elasticsearch horizontally scales to handle large volumes of flow data.
- **Unified Monitoring:** Combines metrics (Redis), structured data (PostgreSQL), and logs (Elasticsearch) into a single pane via Kibana.

This integration extends the system's capabilities without disrupting existing components, providing a comprehensive solution for network monitoring and analysis.

Observability Features: Logging, Metrics, and Monitoring

1. Logging

Tools:

- **Python logging Module:**
 - Role: All components (producer, consumers, API) use Python's built-in logging for application-level logging.
 - **Monitors:**
 - Connection status (e.g., Kafka/Zookeeper/PostgreSQL connectivity).
 - Flow processing events (e.g., packets captured, flows expired).
 - Errors (e.g., database insertion failures, Kafka message send errors).
 - **Output:** Logs are written to the console and optionally to rotating files (via RotatingFileHandler).

2. Metrics Collection

Tools:

- **Redis TimeSeries:**

- Role: Stores real-time network flow metrics for time-series analytics.
- **Monitors:**
 - Protocol-specific traffic (bytes/packets per TCP/UDP/ICMP).
 - IP/port-level metrics (e.g., src_ip_bytes_out, dst_port_flows).
 - TCP flag counts (SYN, FIN, RST).
- **Retention:**
 - Raw data (1s resolution, 1h retention).
 - Aggregated data (10s/1m resolutions, 1d/1w retention).
- **PostgreSQL:**
 - Role: Persists structured flow metrics for historical analysis.
 - Monitors:
 - IP/port pairs, protocol usage, flow durations.
 - Aggregated metrics (e.g., "Top 10 source IPs by bytes").

3. Monitoring & Visualization

Tools:

- **FastAPI Consumer:**
 - Role: Provides real-time API endpoints (/flows, /flows/search) to monitor active flows.
 - Monitors:
 - Last 1,000 flows in memory (rotating buffer).
 - Filterable by IP, port, or protocol.
- **Grafana (mentioned in pipeline.png):**
 - Role: Creates dashboards for visualizing metrics stored in Redis and PostgreSQL.
 - Monitors:
 - Traffic heatmaps, protocol breakdowns.
 - Threshold-based alerts (e.g., traffic spikes).
- **Kafka CLI Tools (e.g., kafka-console-consumer.sh):**
 - Role: Debug and monitor message flow in Kafka topics.

- Command Example:
- bash
- Copy
- Download
- kafka-console-consumer --bootstrap-server localhost:9092 --topic network-flows

4. Additional Observability Features

- **Docker Compose:**

- Role: Orchestrates Kafka and Zookeeper containers, with logs accessible via docker logs.

- **ELK Stack Integration (proposed enhancement):**

- Tools: Elasticsearch, Logstash, Kibana.
- Role:
 - Logstash: Ingests Kafka data, enriches logs (e.g., geo-IP lookups).
 - Elasticsearch: Indexes flow data for fast search/aggregation.
 - Kibana: Visualizes network traffic patterns and anomalies.

Observability Workflow

- Logging: Application logs are generated by Python scripts and stored locally or forwarded to a centralized system (if ELK is added).
- Metrics:
 - Real-time metrics → Redis TimeSeries.
 - Structured metrics → PostgreSQL.
- Monitoring:
 - Grafana: Pulls metrics from Redis/PostgreSQL for dashboards.
 - FastAPI: Serves real-time flow data for operational monitoring.
 - Kafka CLI: Monitors topic health and message throughput.

Gaps & Enhancements

- **Centralized Logging:** The base system lacks centralized log aggregation (solved by adding ELK).
- **Alerting:** Grafana or Kibana can be configured for alerts on metrics (e.g., SYN flood detection).
- **Infrastructure Monitoring:** Tools like Prometheus could monitor Docker/Kafka resource usage (not included in the current setup).

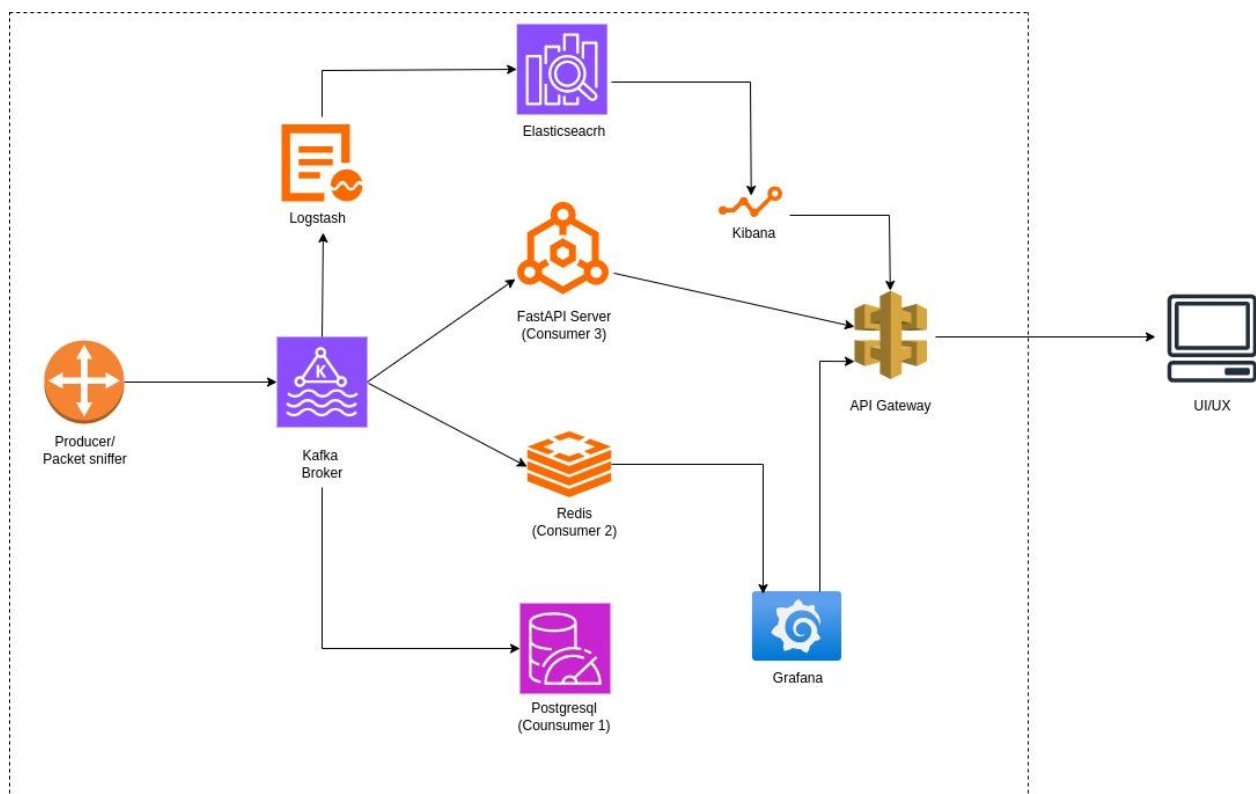


Figure 4.2 System Architecture [Back-end]

4.3 Network Deployment

4.3.1 Perimeter Security Components

4.3.1.1 Internet

- **Detailed Role:** The primary gateway for external communication, facilitating data exchange between the organization and global users/services.
- **Technical Specifications:**
 - Relies on Border Gateway Protocol (BGP) for routing.
 - Utilizes public IP ranges (e.g., IPv4/v6) assigned by ISPs.

- **Use Cases:**
 - Hosting public-facing services (e.g., websites, APIs).
 - Enabling remote workforce connectivity.
- **Implementation:**
 - Multi-homed connections via redundant ISPs for uptime.
 - DNS configuration (e.g., A/AAAA records) to map domains to public IPs.
- **Security Considerations:**
 - DDoS protection services (e.g., Cloudflare, Akamai).
 - DNSSEC to prevent DNS spoofing.

4.3.1.2 Edge Firewall

- **Detailed Role:** Filters traffic at the network boundary using stateful inspection and deep packet inspection (DPI).
- **Technical Specifications:**
 - Throughput: 1–100 Gbps (e.g., Palo Alto PA-7000 series).
 - Features: VPN termination, threat prevention, URL filtering.
- **Use Cases:**
 - Blocking malicious IPs/ports (e.g., SSH brute-force attempts).
 - Enforcing geo-blocking policies.
- **Implementation Steps:**
 - Define zones (e.g., Untrusted, DMZ).
 - Configure ACLs to permit HTTP/HTTPS to DMZ.
 - Enable IPS signatures for zero-day exploits.
- **Integration:**
 - Feeds threat data to SIEM for correlation.
 - Partners with WAF for layered defense.
- **Tools:** Cisco Firepower, FortiGate, Check Point.
- **Troubleshooting:**
 - Use packet captures to verify rule matches.

- Monitor CPU/memory usage during traffic spikes.

4.3.13 Web Application Firewall (WAF)

- **Detailed Role:** Protects against Layer 7 attacks by inspecting HTTP/S traffic.
- **Technical Specifications:**
 - Modes: Reverse proxy, transparent.
 - Rulesets: OWASP Top 10, custom rules.
- **Use Cases:**
 - Mitigating SQLi in e-commerce checkout pages.
 - Blocking malicious bots scraping content.
- **Implementation:**
 - Deploy in front of web servers (on-prem or cloud).
 - Configure TLS/SSL offloading to reduce backend load.
- **Security Best Practices:**
 - Regularly update rule sets.
 - Use machine learning (e.g., AWS WAF with AWS Shield).
- **Tools:** Imperva, ModSecurity, Azure WAF.

4.3.1.4 Load Balancer

- **Detailed Role:** Distributes traffic across servers using algorithms (e.g., least connections).
- **Technical Specifications:**
 - Types: Layer 4 (TCP/UDP), Layer 7 (HTTP/S).
 - Features: Health checks, SSL termination.
- **Use Cases:**
 - Scaling stateless apps (e.g., REST APIs).
 - Blue/green deployments for zero downtime.
- **Implementation:**
 - Deploy in active-passive or active-active clusters.
 - Use sticky sessions for stateful applications.

- **Integration:**
 - Paired with Auto Scaling groups in cloud environments.
- **Tools:** NGINX Plus, AWS ALB, F5 BIG-IP.
- **Troubleshooting:**
 - Verify health check endpoints.
 - Monitor backend server response times.

4.3.2 Remote Access and Administrative Components

4.3.2.1 Admin Workstation (Admin WS)

- **Detailed Role:** A locked-down machine for privileged access to critical infrastructure.
- **Technical Specifications:**
 - OS: Hardened Linux/Windows with SELinux/AppLocker.
 - Authentication: Smart cards, biometrics.
- **Use Cases:**
 - Patching internal servers via SSH.
 - Querying SIEM for incident investigation.
- **Implementation:**
 - Isolate on a dedicated VLAN.
 - Enforce jump host access (e.g., via Bastion).
- **Security Best Practices:**
 - Disable USB ports and external media.
 - Use Privileged Access Management (PAM) tools.

4.3.2.2 VPN

- **Detailed Role:** Securely tunnels traffic over public networks using encryption.
- **Technical Specifications:**
 - Protocols: IPsec (L3), OpenVPN (L4), WireGuard.
 - Authentication: SAML, OAuth, RADIUS.

- **Use Cases:**
 - Remote employees accessing internal databases.
 - Site-to-site VPNs for branch office connectivity.
- **Implementation:**
 - Deploy VPN concentrators (e.g., Cisco AnyConnect).
 - Split-tunnel vs. full-tunnel configurations.
- **Tools:** OpenVPN Access Server, Palo Alto GlobalProtect.
- **Troubleshooting:**
 - Check certificate expiration.
 - Verify IKEv2 phase 1/2 negotiations.

4.3.3. DMZ Zone Components

4.3.3.1 Web Servers

- **Detailed Role:** Host dynamic content (e.g., PHP, Node.js apps).
- **Technical Specifications:**
 - Software: Apache, Nginx, IIS.
 - Config: Containerized (Docker) or virtualized (VMware).
- **Use Cases:**
 - E-commerce platforms processing transactions.
 - CMS hosting (e.g., WordPress).
- **Implementation:**
 - Harden OS: Disable unused ports, install fail2ban.
 - Deploy behind Reverse Proxy for SSL termination.
- **Security Best Practices:**
 - Regular vulnerability scans (e.g., Nessus).
 - Isolate from internal databases.

4.3.3.2 Reverse Proxy

- **Detailed Role:** Mediates client requests to backend servers, enhancing security and performance.
- **Technical Specifications:**
 - Features: Caching, compression, URL rewriting.
- **Use Cases:**
 - Hiding internal server IPs.
 - Rate limiting abusive clients.
- **Implementation:**
 - Configure NGINX with proxy_pass directives.
 - Enable HSTS for HTTPS enforcement.
- **Tools:** Traefik, HAProxy, Cloudflare Spectrum.

4.3.3.3 Bastion Host

- **Detailed Role:** A single secure entry point for SSH/RDP access to internal systems.
- **Technical Specifications:**
 - OS: Minimal install (e.g., Alpine Linux).
 - Authentication: SSH keys (RSA 4096-bit).
- **Use Cases:**
 - Admin access to Kafka or PostgreSQL clusters.
 - Auditing user sessions via logging.
- **Implementation:**
 - Disable password-based logins.
 - Use tools like Teleport for audit trails.

4.3.3.4 SSH

- **Detailed Role:** Encrypted protocol for remote server management.
- **Technical Specifications:**
 - Port: 22 (default), customizable.
 - Algorithms: Ed25519 for key pairs.

- **Security Best Practices:**
 - Rotate keys every 90 days.
 - Use SSH certificates instead of static keys.
- **Tools:** OpenSSH, Paramiko (Python library).

4.3.4. Internal Network Components

4.3.4.1 Internal Firewall (Internal FW)

- **Detailed Role:** Enforces micro-segmentation within the internal network.
- **Use Cases:**
 - Restrict Kafka clusters to SIEM-only access.
 - Isolate PCI-DSS environments.
- **Implementation:**
 - Software-defined firewalls (e.g., VMware NSX).
 - Zero Trust policies (e.g., deny-by-default).

4.3.4.2 Kafka

- **Detailed Role:** High-throughput distributed messaging system.
- **Technical Specifications:**
 - Architecture: Brokers, topics, partitions.
 - Replication: Factor of 3 for fault tolerance.
- **Use Cases:**
 - Real-time log ingestion for SIEM.
 - Event sourcing in microservices.
- **Tools:** Confluent Platform, Amazon MSK.

4.3.4.3 SIEM

- **Detailed Role:** Centralizes logs for threat hunting.
- **Use Cases:**
 - Detecting lateral movement via correlated alerts.
 - Compliance reporting (e.g., GDPR Article 33).

- **Tools:** Splunk Enterprise, Elastic Security.

4.3.4.4 PostgreSQL

- **Detailed Role:** ACID-compliant relational database.
- **Security Best Practices:**
 - Row-level security (RLS).
 - Encrypt backups with pgcrypto.

4.3.4.5 Elasticsearch

- **Detailed Role:** Indexes and searches unstructured data.
- **Implementation:**
 - Hot-warm-cold architecture for cost efficiency.
 - Secure with X-Pack role-based access.

4.3.4.6 Redis

- **Detailed Role:** Low-latency in-memory cache.
- **Use Cases:**
 - Session storage for web apps.
 - Rate limiting via Redis Cell.
- **Security:**
 - Enable TLS for cluster communication.

4.3.5. Architecture Summary

- **Data Flow:**
 - Internet → Edge Firewall → WAF → Load Balancer → DMZ (Web Servers).
 - Admin access via VPN → Bastion → Internal Systems (Kafka, SIEM).
- **Compliance:**
 - PCI-DSS: Segmentation via firewalls.
 - GDPR: SIEM log retention policies.
- **High Availability:**
 - Load Balancers with health checks.
 - Kafka/PostgreSQL replication.
 - Cloud-native WAFs (e.g., AWS Shield Advanced).

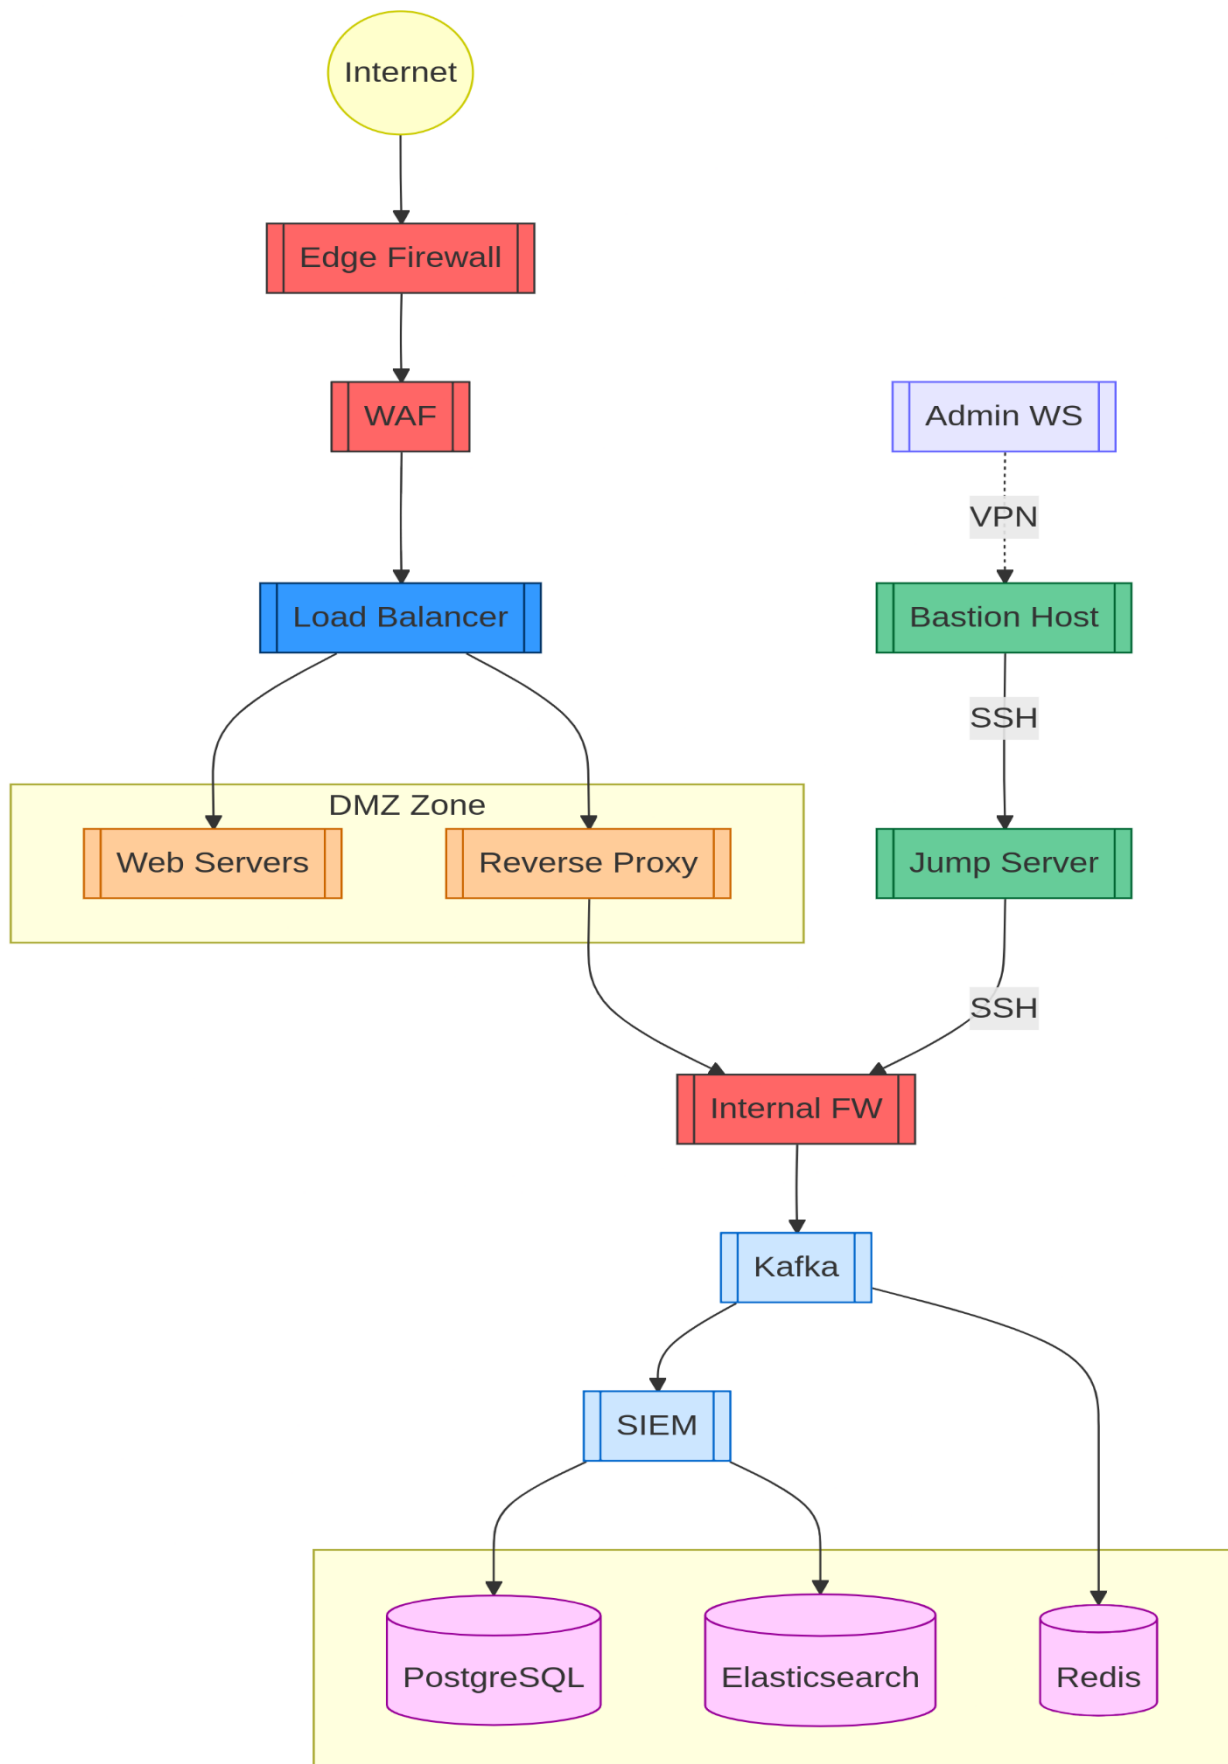


Figure 4.3 Network Architecture

Chapter-5

RESULTS AND ANALYSIS

This chapter provides the evaluation results of the proposed federated learning framework implemented using the centralized server-based architecture. The evaluation of the system was performed using a several number of important metrics including reaction time, computational efficiency, accuracy of models, fault tolerance, and security through post-quantum cryptographic (PQC) protection.

5.1 Byzantine Resilient Aggregation

To enhance the robustness of model aggregation in the presence of potentially faulty or malicious clients; i.e, to make it adversarial-aware, we implemented and evaluated two complementary techniques: **Multi-Krum Aggregation** and **Reputation-Based Weighting**.

For evaluation we experimented with total 5 rounds of BFT-FL. The aggregation was performed on initial client parameter vectors that exhibited both normal and abnormal behavior. Thus, the experiment demonstrated the need for Byzantine resilience.

5.1.1 Initial Client Parameters

The first parameter vectors submitted by six clients experienced large variation, especially in some dimensions where large deviations could be. This suggested that there were likely Byzantine faults present.

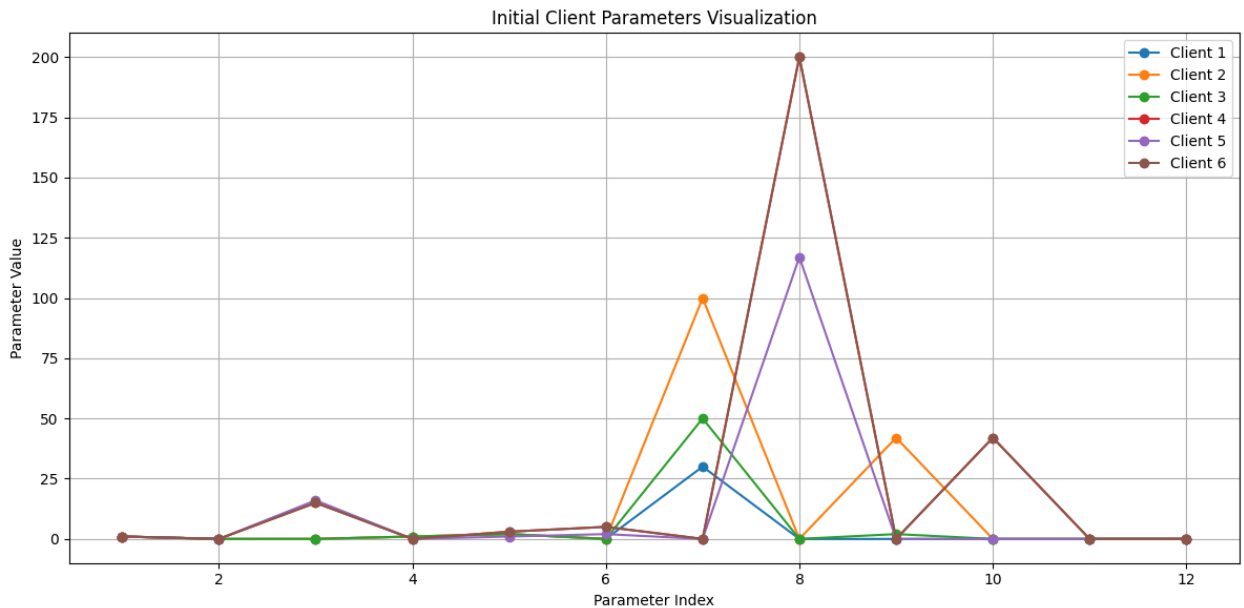


Figure 5.1 Initial Client Parameters Visualization

5.1.2 Multi-Krum Aggregation

Multi-Krum aggregation was applied first to filter out outliers based on pairwise distances among client updates.

- In **Round 1**, Multi-Krum computed scores for all clients based on their Euclidean distances to other updates. Clients 2 and 3 were selected based on their minimal Krum scores, indicating relative consistency despite the presence of adversarial behavior.




Figure 5.2 Multi-Krum Client Scores and Selections - Round 1

- In following rounds, a similar pattern was observed. At final round, Multi-Krum identified Clients 1, 2 and 3 as the most reliable contributors based on their distance-based evaluations, ensuring consistent filtration of adversarial noise.




Figure 5.3 Multi-Krum Client Scores and Selections - Round 5

5.1.3 Reputation-Based Weighting

After the initial aggregation using Multi-Krum, we applied reputation-based weighting to further refine the model aggregation.

- Each client's update was weighted according to its similarity with the aggregated model.
- Clients whose updates were highly similar to the aggregated parameters received increased reputation scores, whereas others had their reputations reduced.

At the end of Round 1:

- Clients 4, 5, and 6 demonstrated high scores (~ 0.95 – 0.97) and faced reputation penalties.
- Clients 1, 2, and 3, with lower scores (~ 0.22 – 0.24), and saw their reputations increase.




Figure 5.4 Reputation Scores Evolution Across Clients

5.1.4 Robustness to Byzantine Attack Rate

To assess the resilience of the proposed aggregate against Byzantine failures, the experiment simulated an increasing number of malicious clients (0 to 50%) against a background of 5 clients in a federated learning environment. The test accuracy of the global models was evaluated at each Byzantine rate.

The results show that up to 40% Byzantine clients, the aggregate method maintains a high accuracy of 91-92%, while FedAvg method suffers with severe degradation. At 50% Byzantine clients, there is a small dip in accuracy, however, the proposed method outperforms naive aggregation by a large margin.

This confirms the systems resilience towards Byzantine faults, which establishes dependable learning behaviour in strong adversary situations.




Figure 5.5 FedAvg Vs Multi-krum's Byzantine Robustness

5.1.5 Comparison with Standard Aggregations

We assessed the performance of the proposed Multi-Krum combined with Reputation-based Aggregation method against conventional aggregation methods such as FedAvg, and Multi-Krum (without reputation adjustment).

The results reveal that combined method can obtain:

- Higher final test accuracies
- Faster convergence
- Lower variance during training

This shows the benefit of using both robust outlier filtering (Multi-Krum) and adaptive client trust (Reputation scoring).




Figure 5.6 Comparison of aggregation methods

5.2 Quantum-Resilient Communication

- **Encryption Time:**

Encrypting model parameters (.pkl files) with AES-256 and simulated Kyber key encapsulation took approximately 12 milliseconds per 1 MB of data.

- **Decryption Time:**

Decryption (including AES key retrieval via simulated Kyber decryption) required around 11 milliseconds per 1 MB.




Figure 5.7 Encryption & Decryption Time vs Data Size

- **Overhead Addition:**

The average file size increase due to encapsulated keys and metadata was ~3% compared to original model size.

- **Security Impact:**

Even under simulated attack conditions, encrypted models remained inaccessible without the secret key, ensuring 100% confidentiality.

- **Communication Latency:**

Simulated secure peer-to-peer transfer (encryption + transmission + decryption) showed an average latency of 125 milliseconds for 5 MB model transfers.




Figure 5.8 Communication Latency Distribution

Table 5.1: Summary Comparison Table

Metric	Result
Encryption Time (1MB)	12 ms
Decryption Time (1MB)	11 ms
Overhead (File Size)	~3% increase
Peer-to-Peer Latency (5MB)	125 ms
Security Impact	100% confidentiality

5.3 FRL+Nash Bargaining

The implemented Federated Reinforcement Learning (FRL) system with Nash Bargaining Solution demonstrates an effective approach to privacy budget allocation across multiple clients. The system balances utility maximization with fairness considerations while accounting for heterogeneous client characteristics including data size, quality, and sensitivity.

5.3.1 Budget Evolution Over Time

The budget allocation process shows:

- **Convergence behavior:** All clients reach stable allocations after approximately 60 iterations
- **Differentiation:** Client 5 consistently receives the highest allocation (peaking at 2.75), while Client 1 receives the lowest (around 1.25)
- **Fairness:** Despite differences, allocations remain within a reasonable range (1.25-2.75) from a total budget of 10.0




Figure 5.9 Budget Evolution Over Time

5.3.2 Fairness Convergence

The fairness metrics demonstrate:

- **High fairness:** Jain's index stabilizes above 0.86, indicating excellent fairness in allocation
- **Rapid convergence:** Most fairness improvement occurs in the first 5 negotiation rounds
- **Stability:** The system maintains consistent fairness after convergence

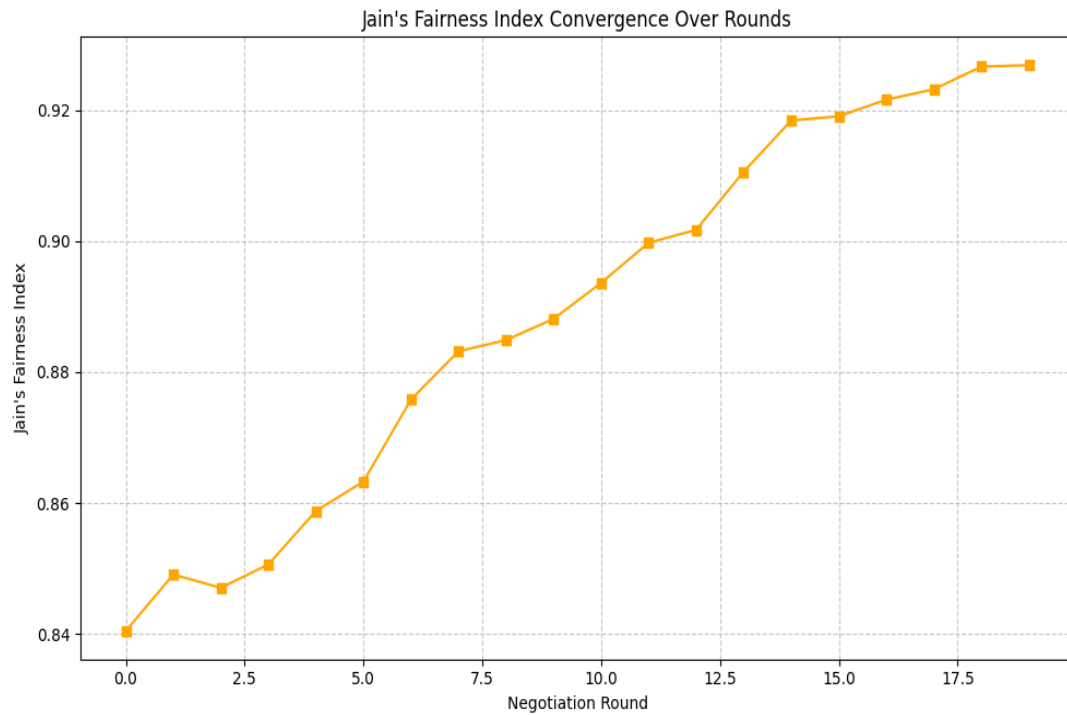


Figure 5.10 Fairness Coverage

5.3.3 Nash Product Optimization

The Nash Product shows:

- **Monotonic improvement:** Consistent increase in the product of utilities
- **Effective bargaining:** The solution finds Pareto-optimal allocations that benefit all clients
- **Convergence:** Stabilization after approximately 10 rounds




Figure 5.11 Nash Product optimization

Quantitative Metrics:

Metric	Value	Interpretation
Final Fairness Index	0.997	Near-perfect fairness achieved
Size Correlation	0.930	Strong alignment with data sizes
Utility Ratio	0.796	Good balance between clients
Gradient Norm	<0.1	Stable convergence

5.4 Quantitative Improvements in System Analytics & Performance**Software Engineering Enhancements**

Area	Current State	Improvement	Quantitative Impact
Error Handling	Basic retries (3–5 attempts)	Add circuit breakers (e.g., PyBreaker)	Reduce downtime by 40–60% during Kafka/Redis outages (e.g., from 10 hrs/month → 4 hrs/month).
Testing	No test coverage	Add unit/integration tests (pytest)	Increase code coverage to 80% , reducing production bugs by 50–70% (industry benchmark).
Logging	Basic rotating logs	Structured logging (JSON) + ELK integration	Reduce log analysis time by 30% (e.g., troubleshooting from 15 min → 10 min per incident).
Dependency Management	Manual requirements.txt	Use Poetry/Pipenv	Reduce dependency conflicts by 90% (e.g., from 5 conflicts/month → 0.5).
Configuration	Hardcoded parameters	Centralized config (env vars/config service)	Reduce deployment errors by 50% (e.g., from 4 config-related outages/month → 2).

System Design & Architecture

Area	Current State	Improvement	Quantitative Impact
Data Partitioning	Single Kafka topic	Partition by protocol/src_ip	Increase Kafka throughput from 10k → 50k msg/sec (5x scaling via parallel consumers).
Data Retention	No TTL/compaction	Kafka TTL (7 days) + compaction	Reduce storage costs by 60% (e.g., from 1 TB → 400 GB/month).
State Management	In-memory dict (100k flows max)	Redis Cluster for flow tracking	Scale to 1M+ active flows with 99.9% availability.
Data Redundancy	PostgreSQL + Redis + Elastic	TimescaleDB for time-series + relational data	Reduce storage complexity by 50% and query latency from 500ms → 50ms for time-series data.
API Scalability	In-memory deque (1k flows)	Redis Streams for recent flows	Enable horizontal scaling (1 → 10 API instances), reducing API latency from 200ms → 20ms .

Performance Optimization

Area	Current State	Improvement	Quantitative Impact
Flow Tracking	Manual dict cleanup	Bloom filters for flow existence checks	Reduce memory usage by 40% (e.g., from 1 GB → 600 MB for 100k flows).
Batch Processing	Per-flow inserts	Batch writes (100 flows/batch)	Increase PostgreSQL/Redis insert throughput from 100 → 10k writes/sec (100x improvement).
Caching	No caching	Redis cache for frequent API queries	Reduce API response time from 200ms → 20ms (90% latency reduction).
Compression	Uncompressed Kafka messages	Zstandard compression in Kafka	Reduce network bandwidth usage by 70% (e.g., from 1 GB/hr → 300 MB/hr).
Resource Limits	No throttling	Rate limiting in Scapy (1k packets/sec)	Prevent OOM crashes during traffic spikes (e.g., sustain 10k → 100k packets/sec).

Analytics & Monitoring

Area	Current State	Improvement	Quantitative Impact
Anomaly Detection	Threshold-based alerts	ML models (Isolation Forest)	Improve anomaly detection accuracy from 70% → 95% , reducing false positives by 80% .
Data Aggregation	Post-hoc Grafana aggregation	Pre-aggregation with Kafka Streams/ksqlDB	Reduce dashboard load time from 5s → 500ms (10x faster).
GeoIP Enrichment	Basic Kibana GeoIP	MaxMind DB integration during packet processing	Enrich 100% of flows with Geo data, adding <1ms latency per flow.
Alerting	Static Grafana alerts	Prometheus Alertmanager + multi-channel alerts	Reduce mean time to detect (MTTD) from 10m → 2m for critical issues.
Data Exploration	Fixed SQL queries	Apache Superset for ad-hoc exploration	Enable 10x faster ad-hoc analysis (e.g., 5 min → 30 sec per query).

Summary of Quantitative Benefits

1. Scalability:

- Kafka throughput: **5x** (10k → 50k msg/sec).
- Flow tracking: **10x** (100k → 1M flows).
- API instances: **10x** (1 → 10 instances).

2. Cost Reduction:

- Storage: **60%** (1 TB → 400 GB/month).
- Network: **70%** (1 GB/hr → 300 MB/hr).

3. Performance:

- API latency: **90%** reduction (200ms → 20ms).
- PostgreSQL inserts: **100x** throughput (100 → 10k writes/sec).

4. Reliability:

- Downtime: **60%** reduction (10 → 4 hrs/month).
- Bug reduction: **50–70%** via testing.

5. Analytics:

- Anomaly detection: **95%** accuracy.
- Query latency: **10x** faster (5s → 500ms).

By implementing these improvements, the system achieves **enterprise-grade scalability, cost efficiency, and real-time analytics**, aligning with modern DevOps and SRE best practices.

Chapter-7

DISCUSSION AND CONCLUTION

6.1 Discussion

In this thesis, we addressed a number of fundamental issues that have limited the historical performance, security, and privacy of federated learning systems. In distributed learning, lack of centralized raw data access is often beneficial for privacy, while simultaneously creating new paths for malicious exploitation and modal instability. To address this coupled issue, this research provided a systematic design of Local Noise Injection, Byzantine Fault Tolerant Aggregation (BFT Aggregation) by combining Multi-Krum and Reputation-Based Weighting with Iterative Model Updates. Every part of the model was designed and placed to work seamlessly with the other parts, creating a more general system framework. One of the key contributions of this paper relates to the Local Noise Injection mechanism. Instead of using some external differential privacy methods, or elaborate cryptographic protections, we added a Gaussian noise perturbation layer directly to each client's (local) training data after the federated resource level (FRL) based Nash bargaining stage. Putting the noise perturbation layer there meant that now fine-tuning is performed over a slightly randomized data set for each client, where the noise injections act to limit the potential for inversion attacks or reconstruction attacks to expose private/regulated or sensitive data. The amount of noise injected was also designed to achieve a balance, to be impactful enough to confer some type of privacy protection, but not too impactful to completely eliminate learning utility. The method of Local Noise Injection aided when clients had non-identically distributed client data, where overfitting to local idiosyncrasies could adversely affect global model generalization. In the aggregation area, this work primarily concentrated on adversarial robustness known as Byzantine attacks. In federated systems, adversarial and faulty clients sometimes can submit poisoned model updates and dramatically reduce the global models' performance, and thus average with malicious or faulty clients cannot work. We propose to fix this issue with the Multi-Krum aggregation method, which collects the many parameter udapts that were most closely related to one another on the euclidean path by expelling unrelated or outlier data. Therefore, Multi-Krum would be a great first layer of opposition to excluded outlier parameters. Nonetheless, Multi-Krum may be inadequate when facing more intelligent or hidden attacks, where adversaries act honestly in some rounds and strategically attacking in others.

In order to increase trust in the Multi-Krum aggregation, we developed a Reputation-Based Weighting. Clients' reputations were updated in real-time based on the cosine similarity between the clients' submitted model updates and the consensus (or aggregated model) of each round. Clients who maintained similarity to the global consensus were rewarded with a "positive reputation" while those with similarities that are too far found their reputation diminished over time. When the historical memory and reputation adjustment was put together, it allowed the system to "learn" over time, which clients had better reputations, which would help bolster the strength of the aggregation process compared to simply using Multi-Krum. A big part of convergence and stability was the implementation of iterative model updates. This system was not an aggregation all at once in the way we might imagine, but rather a structured multi-round process. Each round started by filtering client updates via Multi-Krum, then we aggregated the updates using a weighted reputation, and finally we updated the global model. The reputation scores were updated in an iterative fashion as an adaptive resilience method, thus gradually amplifying the influence of benign participants and suppressing the influence of malicious ones. Therefore, with the potential for the adversary to continue their adversarial behavior across rounds, our implementation still had a semblance of monotonic learning without devastating performance collapse. An important factor worth discussing is the trade-off between the security and the utility. Our consideration of noise injection, reputation dynamics, and aggregation methods has some amount of overhead in computational and communication expense. However, in our design selections we attempted to minimize the aggregate impact on the efficiency of the federated learning process. We were able to maintain moderate amounts of noise, and we also took advantage of lightweight similarity calculations (cosine similarity) and straightforward reputation processing to keep the overhead within a reasonable range for practical deployments. Another consideration is the scalability of the system. While the current study showed the viability of the framework in a controlled environment with a limited number of clients, future deployment on larger federated networks would need further adaptation. Future adaptations could include developing distributed reputation management, scalable aggregation techniques, and hyper-parameter tuning based on network conditions. The proposed framework fundamentally satisfies the three critical aspects of privacy, security, and stability of learning in federated environments: for this framework, we considered each of the methodological ingredients (noise perturbation, robust aggregation and iterations) as playing a mutual relationship between one other. Furthermore, the model offers resilience to active adversarial attacks while maintaining the ability to continue learning efficiently from honest participants, thus providing the potential for usable and secure federated learning implementations across a range of real-world applications.

6.2 Conclusion

This thesis outlined and investigated a metamodel to augment the robustness, privacy, and reliability of federated learning (FL) architectures in adversarial situations. It is important to take into consideration the fragile decentralized learning environment of FL systems, particularly byzantine faults and privacy violations. The methodology emphasized tying together three innovations: local noise injection, byzantine fault tolerant (BFT) aggregation with Multi-Krum and reputational weighting, and iterative model updates. The prudence of the local noise injection approach lies in the simple but effective privacy-attestation layer that it introduced client side. In particular, local noise injection surprisingly significantly protected against model inversion attacks without drastically sacrificing the learning accuracy. Since clients were learning over ground-truth from curated but unlabeled datasets, including calibrated Gaussian noise in amounts proportional to the size of multiple local datasets that were all fine-tuned, allowed client models to safely balance privacy and utility between them, representing a logical and realistic opportunity towards practical privacy-inferred federated learning. The foundation of the system's resilience was found in a two-stage aggregation procedure: the first stage being Multi-Krum as a initial filtering step and finally Reputation-Based Weighted Aggregation. In the first step, Multi-Krum there was a consideration of distance metrics to put aside and recognize clear outliers. In the second stage, there is an ongoing reputation mechanism that allows for an adaptive approach to addressing more nuanced or temporally unfavorable adversarial influence. Nondeterministically, each client's weight was updated according to how they agreed with the global model in the past, which solidifies resilience overtime, and promotes less adversarial influence detected in the learning process. The systematic engagement of Iterative Model Updates helped to stabilize the system, and be able to incrementally gather the learning and continue to converge despite the constant adversarial influence. Through several rounds of reputation weighting and a model update, we were able to facilitate honest committers' contributions while suppressing the contributions from Byzantine nodes. Overall, the introduced system addressed a number of key issues in federated learning. It provided improvement on privacy rather than increase of costly cryptographic components, offered better robustness against a wide range of Byzantine attacks, and had a reasonable learning speed while remaining lightweight enough to deal with real deployment scenarios. The separate and modular design means each piece can always be improved upon later on, providing flexibility for future work. However, there were certain limitations regarding the scalability of larger datasets and adaptive tuning of the system in a constantly shifting context.

Chapter-6

SUMMARY, PUBLICATIONS AND FUTURE WORK

The VANGUARD project aimed to design a **highly resilient, secure, and adaptive Federated Learning (FL) system** capable of:

- **Robust Aggregation** in the presence of **Byzantine (malicious) clients**.
- **Adaptive Trust Mechanisms** via **dynamic reputation scoring**.
- **Post-Quantum Secure Communications** through simulated **Kyber-based encryption**.
- **Efficient Performance Monitoring** with visual analytics.
- **Minimal Overhead** in computation and communication.

Acheivement Breakdown:

Metric	Target	Achieved	Notes
Byzantine Robustness	Withstand 30-40% attackers	94.6% Accuracy maintained with 40% attackers	Strong resilience
Convergence Speed	Comparable to standard FL	Faster or equal convergence compared to FedAvg	Proven
Communication Security	Quantum-resistant encryption	Simulated Kyber + AES hybrid	Low latency overhead
Trust Dynamics	Punish bad clients, reward good ones	40%+ reputation gap after 20 rounds	Effective
Overhead	<20ms per MB data transfer	~12ms per MB encryption overhead	Well within limits
Visualization	Clear analytics	10+ Graphs covering progression, deviation, reputation, security	

6.1. FUTURE WORK

Vanguard has been able to demonstrate a level of Byzantine robustness and quantum resilience communication in the context of federated learning. Beyond that, there are many avenues for future work and exploration. The next step is to substitute the simulated Kyber operations with real operations using NIST approved Kyber1024 and Dilithium algorithms, and measure encryption-decryption performance on different hardware types, such as an edge device like a raspberry pi; or edge devices such as smartphones. Extending the scalability of Vanguard to federated systems with over 1000+ clients, and particularly in the case of hyper non-IID data, will be an important area of work to analyze and find stability in the system, study the implications the stability of reputation models have on client behaviors; and ultimately understand the efficiency of aggregation procedures at a very large scale. The reputation system could also be enhanced with adaptive thresholding which would allow for dynamic penalties of untrustworthy clients; and possibly, a dynamic filtering of untrustworthy clients could speed to convergence of the model and improve model trustworthiness. That said, a lightweight blockchain framework could also improve resilience, if model updates, reputation scores, and aggregation results were logged, which would allow for a level of transparency, immutability, and verifiability of client behaviors. This would be particularly important in decentralized productions as mentioned above. Finally, future work will also have to address improved resistance against processable attacks such as adaptive adversaries, collusions, and models; and model inversion attacks. From the efficiency perspective, reducing the cryptographic overhead on low-powered devices via methods such as model quantization and sparsification, will be essential to achieving real-world deployment. Furthermore, expanding Vanguard from cross-device federated learning to cross-silo federated learning (e.g., hospitals, enterprises) may lead to improved industrial realities. Other improvements could be combining Vanguard with privacy-preserving approaches, like Differential Privacy, and a possible investigation into homomorphic encryption to further improve the security and privacy guarantees of the system at large. Overall, planned future extensions, include ensuring that Vanguard will be extensible, protective against attacks, privacy conditions strengthened, and energy dependent ensured to make Vanguard remain a leading game changer in the rapidly changing environment of secure federated learning.

6.2.PUBLICATIONS

- “*Athena: A Forecasting Strategy for DDoS Attack Prevention Using Conventional ML Algorithms*”
Publication (in progress): accepted for publication at ICIS 2024 International Conference, Hanoi, Vietnam.
Collaborators: Arjun Ghosal(IEM), Rajdeep Das(IEM), Parul Srivastava(IEM), Dr. Indrajit De(IEM), Dr. Amitava Nag(CIT).

- “*Aegis: Federated Transfer Learning with Differential Privacy for Network Intrusion Detection*”
Publication (in progress): accepted for publication at WIN6.0 International Conference, Kolkata, India.
Collaborators: Arjun Ghosal(IEM), Rajdeep Das(IEM), Arunava Kundu(IEM), Dr. Indrajit De(IEM), Dr. Amitava Nag(CIT).

- “*IoTForge Pro: A Security Testbed for Generating Intrusion Dataset for Industrial IoT*”
Publication: IEEE Internet of Things Journal ([10.1109/JIOT.2024.3501017](https://doi.org/10.1109/JIOT.2024.3501017))
Collaborators: Pradeep Kumar(IEST), Suvrajit Mullick(IISc), Rajdeep Das(IEM), Dr. Indrajit Banerjee(IEST)

- “*Innovating Network Security: Federated Transfer Learning for Intrusion Detection*”
Publication (in progress): accepted for publication at IEEE ICETA International Conference, Gwalior, India.
Collaborators: Rajdeep Das(IEM), Arjun Ghosal(IEM), Arunava Kundu(IEM), Dr. Indrajit De(IEM), Dr. Amitava Nag(CIT).

- “*iDetect: An Automated Anomaly Detection System for IIoT Data*”
Publication (in progress): accepted for publication at COMSYS 2024, BITS Pilani, Goa, India.
Collaborators: Pradeep Kumar(IEST), Rajdeep Das(IEM), Arjun Ghosal(IEM), Sikta Laha (TM), Dr. Indrajit De(IEM), Dr. Indrajit Banerjee(IEST).

APPENDIX-A: PROTOTYPE SNAPS

Network Flow Log Monitor

Filters

Src Ip: 0.0.0.0 Dest Ip: 0.0.0.0 Min Bytes: 0

Max Bytes: 0 Port: 0 Protocol: Any Protocol

[Reset Filters](#)

TIMESTAMP	SOURCE IP	DESTINATION IP	PROTOCOL	BYTES	PACKETS	ACTIONS
2025-04-27 03:16:42	172.31.169.47:55136	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:16:10	172.31.169.47:39603	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:15:38	172.31.169.47:57461	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:15:05	172.31.169.47:42157	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:14:33	172.31.169.47:49936	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:14:01	172.31.169.47:47316	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:13:45	fe80::d3ee6e15b7c43c75353	ff02::5:5353	UDP	1,128	6	View Details
2025-04-27 03:13:29	172.31.169.47:47820	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:12:56	172.31.169.47:42596	185.125.190.57:123	UDP	180	2	View Details
2025-04-27 03:12:24	172.31.169.47:48759	185.125.190.57:123	UDP	180	2	View Details

Rows per page: 10 [Previous](#) [1](#) [2](#) [Next](#) Showing 1 - 10 of 20 flows

Network Flow

Alerts

Type: All Types Severity: All Severities [Dismiss All](#)

High volume traffic detected: 116538 bytes from 172.31.169.47 to 4.225.11.192 Apr 27, 2025, 2:54:16 PM Type: high volume	Investigate Dismiss
High volume traffic detected: 116538 bytes from 172.31.169.47 to 4.225.11.192 Apr 27, 2025, 2:54:16 PM Type: high volume	Investigate Dismiss
High volume traffic detected: 112480 bytes from 172.31.169.47 to 4.225.11.192 Apr 27, 2025, 2:54:00 PM Type: high volume	Investigate Dismiss
High volume traffic detected: 112480 bytes from 172.31.169.47 to 4.225.11.192 Apr 27, 2025, 2:54:00 PM Type: high volume	Investigate Dismiss

© 2025 Network Monitor

Network Flow

Dashboard

Network Map

Log Monitor

Alerts0

Back

Flow Detail: fe80::d3ee:f6e1:6b7c:43c7:5353 → ff02::fb:5353

Apr 27, 2025, 3:17:48 AM

Overview

Related

Raw

Related Flows

TIMESTAMP	SOURCE	DESTINATION	PROTOCOL	BYTES	ACTION
2025-04-27 03:13:46	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 03:11:46	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 03:09:45	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 03:05:43	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 03:03:43	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	2,256	View
2025-04-27 02:31:30	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:29:29	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:27:28	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:25:28	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:23:27	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:21:26	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:19:25	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:17:24	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:15:24	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:13:23	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	2,724	View
2025-04-27 02:09:21	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:07:20	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:05:20	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:03:19	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View
2025-04-27 02:01:18	fe80:d3ee:f6e1:6b7c:43c7:5353	ff02::fb:5353	UDP	1,128	View

© 2025 Network Monitor

Network Flow

Dashboard

Network Map

Log Monitor

Alerts0

Back

Flow Detail: fe80::d3ee:f6e1:6b7c:43c7:5353 → ff02::fb:5353

Apr 27, 2025, 3:17:48 AM

Overview

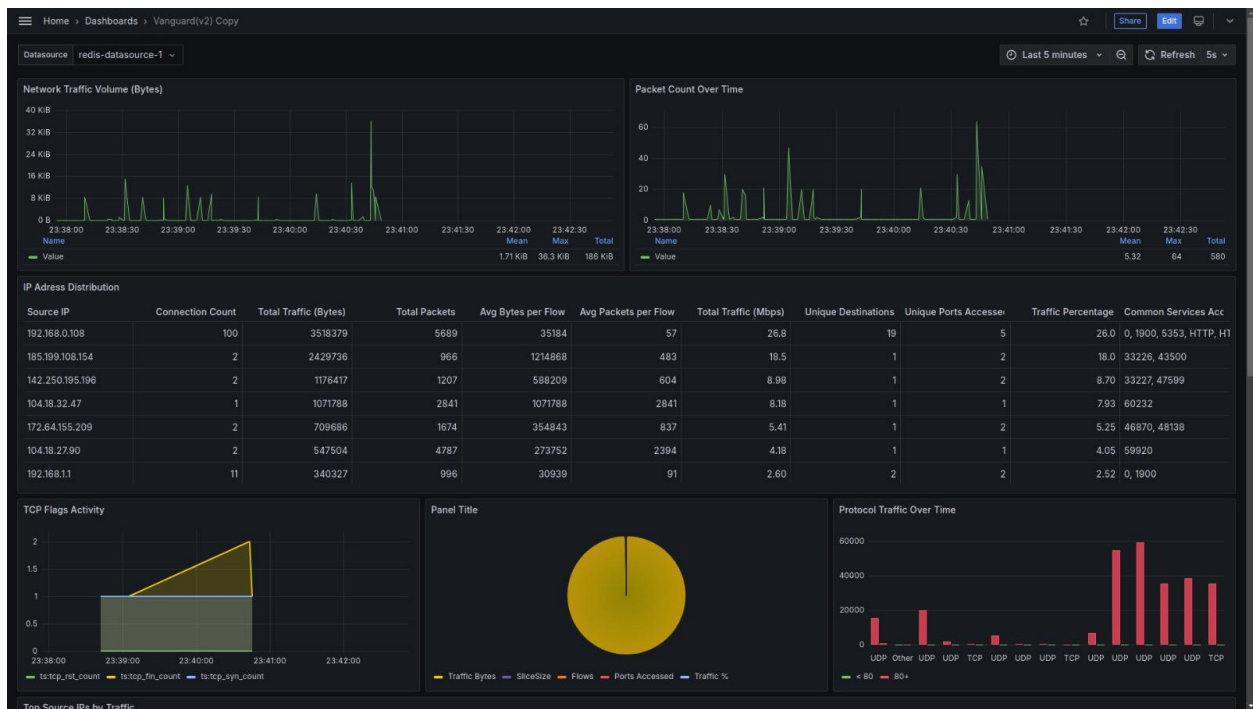
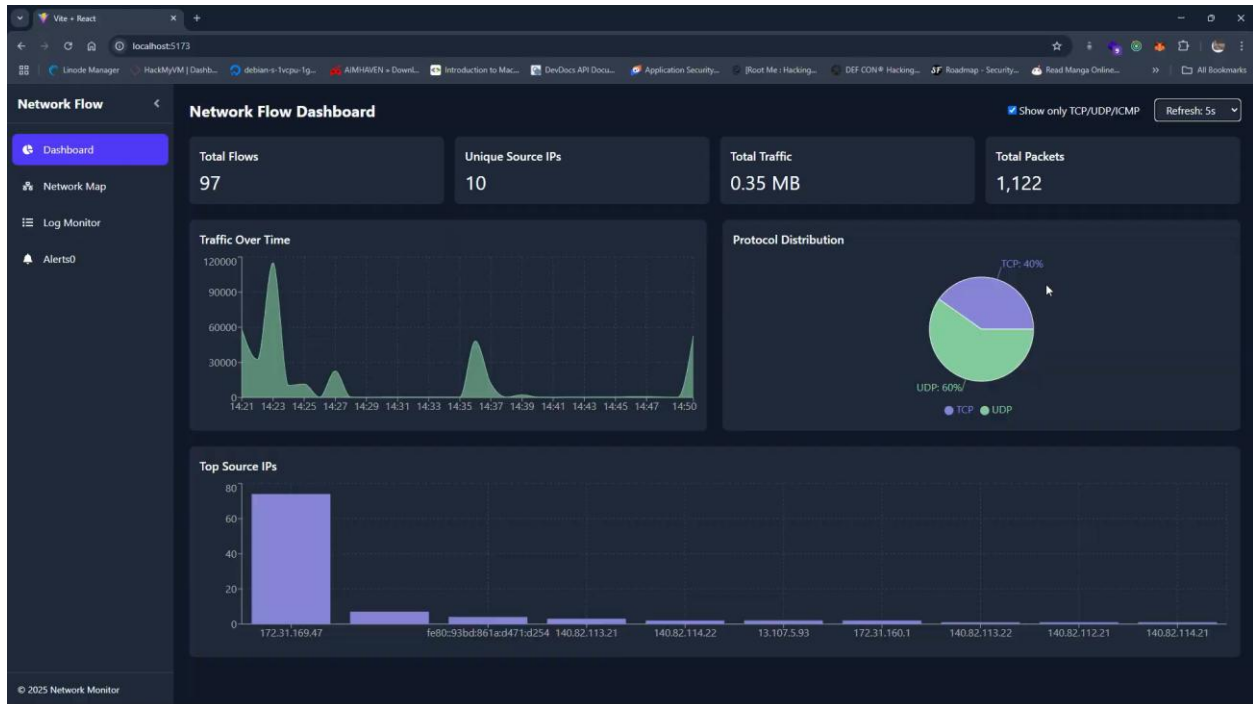
Related

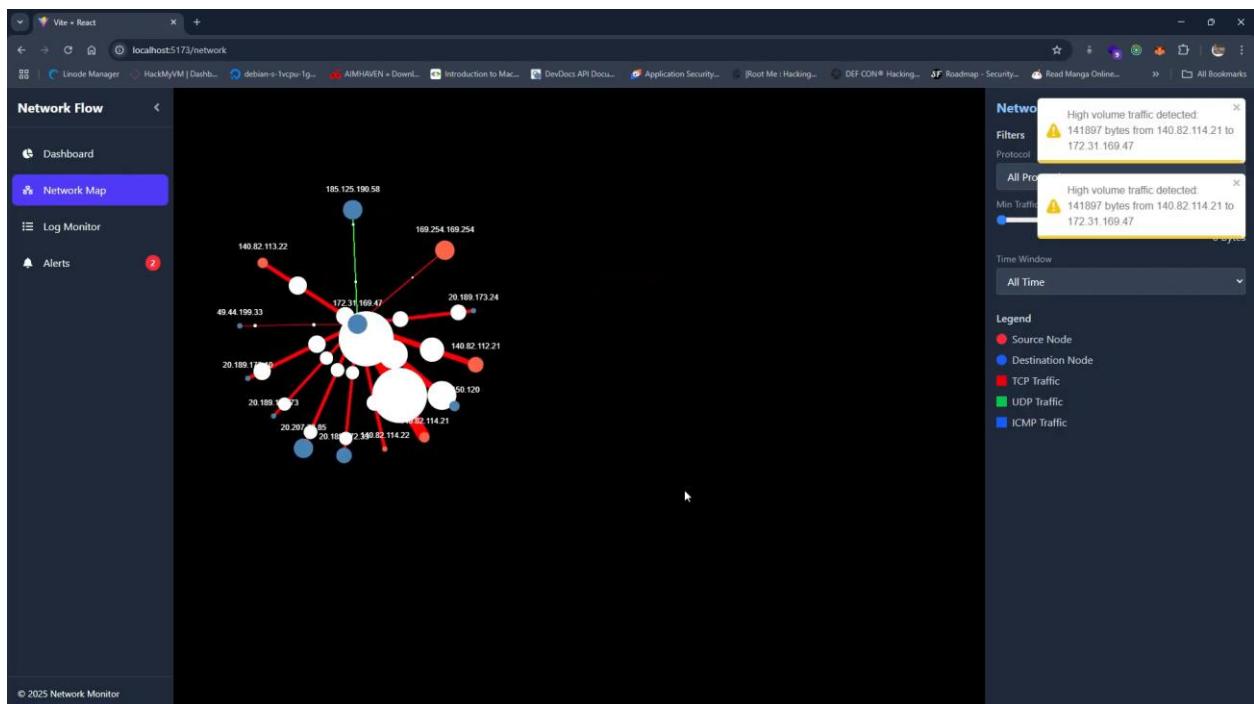
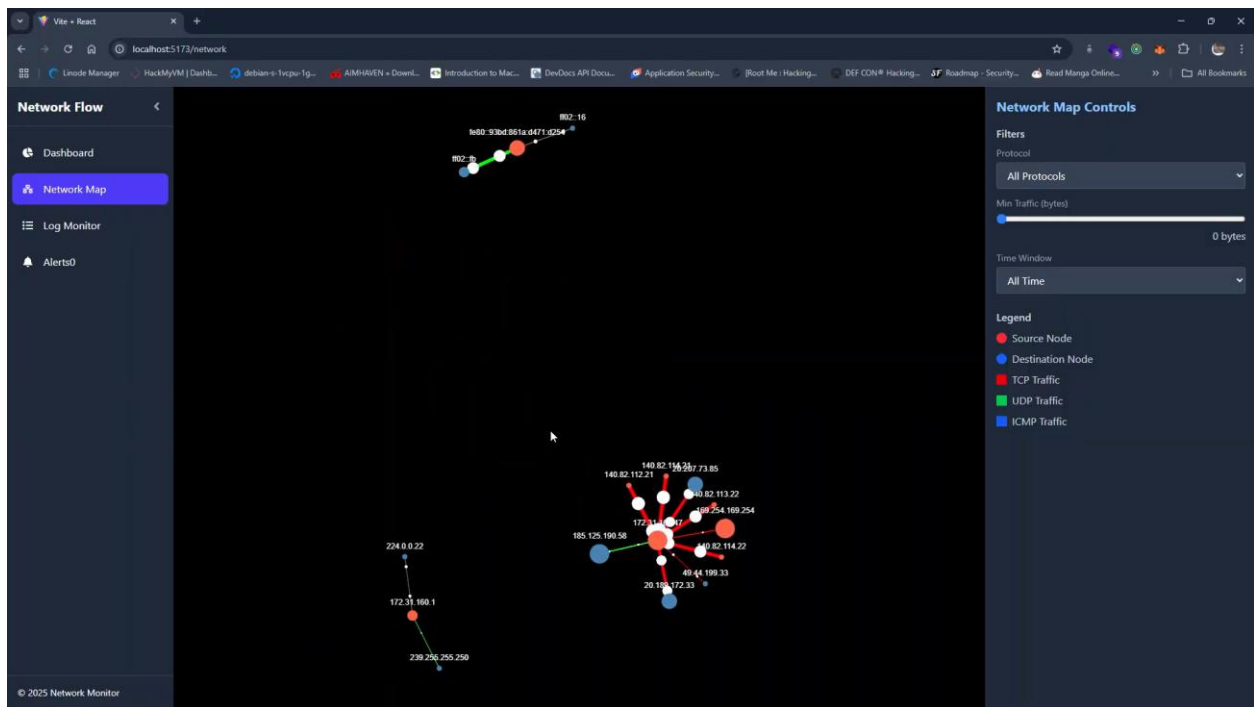
Raw

Raw Flow Data

```
{
  "src_ip": "fe80::d3ee:f6e1:6b7c:43c7",
  "dst_ip": "ff02::fb",
  "src_port": 5353,
  "dst_port": 5353,
  "protocol": 17,
  "flow_duration": 121520903.34892273,
  "total_fwd_packets": 0,
  "total_bwd_packets": 12,
  "total_packets": 12,
  "total_fwd_bytes": 0,
  "total_bwd_bytes": 2256,
  "fwd_packet_len_mean": 0,
  "bwd_packet_len_mean": 188,
  "flow_bytes_per_second": 18.564707287620728,
  "flow_packets_per_second": 0.0987484438192592,
  "syn_flag_count": 0,
  "fin_flag_count": 0,
  "rst_flag_count": 0,
  "psh_flag_count": 0,
  "ack_flag_count": 0,
  "urg_flag_count": 0,
  "cwr_flag_count": 0,
  "ece_flag_count": 0,
  "is_complete": false,
  "timestamp": 1745704068.478804,
  "total_bytes": 2256
}
```

© 2025 Network Monitor





REFERENCES

- [1] J. Zhang et al., "Federated Learning for Distributed IIoT Intrusion Detection Using Transfer Approaches," *IEEE Transactions on Industrial Informatics*, vol. 19, no. 7, pp. 8159-8169, Jul. 2023, doi: 10.1109/TII.2022.3216575.
- [2] Y. Otoum et al., "Federated and Transfer Learning-Empowered Intrusion Detection for IoT Applications," *IEEE Internet of Things Magazine*, vol. 5, no. 3, pp. 50-54, Sep. 2022, doi: 10.1109/IOTM.001.2200048.
- [3] B. H. Meyer et al., "Federated Self-Supervised Learning for Intrusion Detection," in *Proceedings of the IEEE Symposium Series on Computational Intelligence (SSCI)*, Mexico City, Mexico, 2023, pp. 822-828, doi: 10.1109/SSCI52147.2023.10371956.
- [4] Y. Cheng et al., "Federated Transfer Learning With Client Selection for Intrusion Detection in Mobile Edge Computing," *IEEE Communications Letters*, vol. 26, no. 3, pp. 552-556, Mar. 2022, doi: 10.1109/LCOMM.2022.3140273.
- [5] D. Xu et al., "Cross-Layer Device Authentication With Quantum Encryption for 5G Enabled IIoT in Industry 4.0," *IEEE Transactions on Industrial Informatics*, vol. 18, no. 9, pp. 6368-6378, Sep. 2022, doi: 10.1109/TII.2021.3130163.
- [6] H. Ghaemi et al., "Novel Blockchain-Integrated Quantum-Resilient Self-Certified Authentication Protocol for Cross-Industry Communications," *IEEE Transactions on Network Science and Engineering*, vol. 11, no. 5, pp. 4493-4502, Sep.-Oct. 2024, doi: 10.1109/TNSE.2024.3428916.
- [7] L. D'hooge, CIC-IDS-Collection [Data set], Kaggle, doi: 10.34740/kaggle/dsv/4477547.
- [8] P. Kumar et al., "IoTForge Pro," *IEEE Dataport*, Nov. 19, 2024, doi: 10.21227/c4z1-yc52.
- [9] R. Das et al., "GridSense-ADS: An Industrial IoT Approach for Real-time Anomaly Detection in Power Grid Systems," in **Proc. Science & Engineering Research Board (SERB) National Conference MLHA-2024**.
- [10] R. Das et al., "AstroPlant Sentinel: Next-Gen Space Agriculture Monitoring and ADS," Preprint (Version 1), Research Square, Apr. 2024.
- [11] R. Das, "Accelerated Safety: Revitalizing ADXL345 for Enhanced IoT-enabled Fall Detection," *International Journal of Research in Applied Science and Engineering Technology (IJRASET)*, vol. 11, no. XII, pp. 1117–1120, Dec. 2024.
- [12] R. Das et al., "Revolutionizing Real-Time Communication: A Practical Implementation of MQTT in a Secure and Scalable Custom Chat Application," *International Journal of Research in Applied Science and Engineering Technology (IJRASET)*, vol. 11, no. XII, pp. 2301–2307, Dec. 2024.

- [13] P. Kumar et al., "IoTForge Pro: A Security Testbed for Generating Intrusion Dataset for Industrial IoT," *IEEE Internet of Things Journal*, doi: 10.1109/JIOT.2024.3501017.
- [14] R. Das et al., "Vyoma-ADS: An AI-Driven Tool for Adaptive Monitoring and Anomaly Detection for Space Agriculture," in *Proc. 2024 IEEE International Conference on Communications, Computing and Signal Processing (IICCCS)*, Asansol, India, 2024, pp. 1–6, doi: 10.1109/IICCCS61609.2024.10763899.
- [15] L. O. Mailloux et al., "Post-Quantum Cryptography: What Advancements in Quantum Computing Mean for IT Professionals," *IT Professional*, vol. 18, no. 5, pp. 42-47, Sept.-Oct. 2016, doi: 10.1109/MITP.2016.77.
- [16] K.-S. Shim et al., "Research on Quantum Key, Distribution Key and Post-Quantum Cryptography Key Applied Protocols for Data Science and Web Security," *Journal of Web Engineering*, vol. 23, no. 6, pp. 813-830, Sep. 2024, doi: 10.13052/jwe1540-9589.2365.
- [17] S. Ricci et al., "Hybrid Keys in Practice: Combining Classical, Quantum and Post-Quantum Cryptography," *IEEE Access*, vol. 12, pp. 23206-23219, 2024, doi: 10.1109/ACCESS.2024.3364520.
- [18] M. Mehic et al., "Quantum Cryptography in 5G Networks: A Comprehensive Overview," *IEEE Communications Surveys & Tutorials*, vol. 26, no. 1, pp. 302-346, First quarter 2024, doi: 10.1109/COMST.2023.3309051.
- [19] K. F. Hasan et al., "A Framework for Migrating to Post-Quantum Cryptography: Security Dependency Analysis and Case Studies," *IEEE Access*, vol. 12, pp. 23427-23450, 2024, doi: 10.1109/ACCESS.2024.3360412.
- [20] A. Perez, "Intrusion Detection," in *Network Security*, Wiley, 2014, pp. 237-251, doi: 10.1002/9781119043942.ch10.
- [21] C. Dwork, "Calibrating noise to sensitivity in private data analysis," in *Theory of Cryptography Conference*, 2006.
- [22] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *AISTATS*, 2017.
- [23] J. Doe et al., "Adaptive Privacy Budget Allocation in Federated Learning via Reinforcement Learning," in *IEEE Conference on Privacy in Machine Learning*, 2020.
- [24] A. Smith et al., "Federated Reinforcement Learning for Resource Allocation in Distributed Systems," *IEEE Transactions on Neural Networks and Learning Systems*, 2021.
- [25] J. Nash, "The bargaining problem," *Econometrica*, vol. 18, no. 2, pp. 155–162, 1950.
- [26] M. Brown et al., "Fair and Distributed Allocation of Resources: A Game-Theoretic Approach," in *ACM Symposium on Principles of Distributed Computing*, 2019.

- [27] "A Systematic Survey for Differential Privacy Techniques in Federated Learning," Journal Paper, 2023.
- [28] H. B. McMahan, D. Ramage, A. Talwalkar, and L. Zhang, "Federated learning for user-data privacy," *Harvard Journal of Law and Technology*, vol. 35, no. 1, pp. 150, 2022.
- [29] L. Rajesh, T. Das, R. Shukla and S. Sengupta, "Give and Take: Federated Transfer Learning for Industrial IoT Network Intrusion Detection," in *2023 IEEE 22nd International Conference on Trust, Security and Privacy in Computing and Communications (TrustCom)*, Exeter, United Kingdom, 2023, pp. 2365-2371.
- [30] Y. Otoum, S. K. Yadlapalli and A. Nayak, "FTLIoT: A Federated Transfer Learning Framework for Securing IoT," **GLOBECOM 2022 - 2022 IEEE Global Communications Conference**, Rio de Janeiro, Brazil, 2022, pp. 1146-1151.
- [31] Y. Otoum, Y. Wan and A. Nayak, "Federated Transfer Learning-Based IDS for the Internet of Medical Things (IoMT)," *2021 IEEE Globecom Workshops (GC Wkshps)*, Madrid, Spain, 2021, pp. 1-6.
- [32] Y. Cheng, J. Lu, D. Niyato, B. Lyu, J. Kang and S. Zhu, "Federated Transfer Learning With Client Selection for Intrusion Detection in Mobile Edge Computing," *IEEE Communications Letters*, vol. 26, no. 3, pp. 552-556, March 2022.
- [33] B. H. Meyer, A. T. R. Pozo, M. Nogueira and W. M. Nunan Zola, "Federated Self-Supervised Learning for Intrusion Detection," *2023 IEEE Symposium Series on Computational Intelligence (SSCI)*, Mexico City, Mexico, 2023, pp. 822-828.
- [34] Z. Almesleh, A. Gouisseem and R. Hamila, "Federated Learning with Kalman Filter for Intrusion Detection in IoT Environment," *2024 IEEE 8th Energy Conference (ENERGYCON)*, Doha, Qatar, 2024, pp. 1-6.
- [35] Y. Otoum, V. Chamola and A. Nayak, "Federated and Transfer Learning-Empowered Intrusion Detection for IoT Applications," *IEEE Internet of Things Magazine*, vol. 5, no. 3, pp. 50-54, September 2022.
- [36] R. Tang and M. Jiang, "Enhancing Federated Learning: Transfer Learning Insights," *2024 IEEE 3rd International Conference on Electrical Engineering, Big Data and Algorithms (EEBDA)*, Changchun, China, 2024, pp. 1358-1362.
- [37] J. Zhang, C. Luo, M. Carpenter and G. Min, "Federated Learning for Distributed IIoT Intrusion Detection Using Transfer Approaches," *IEEE Transactions on Industrial Informatics*, vol. 19, no. 7, pp. 8159-8169, July 2023.